

# A Discussion of Selected Vienna-Libraries for Computational Science

Karl Rupp<sup>1</sup>, Florian Rudolf<sup>2</sup>, Josef Weinbub<sup>2</sup>



<sup>1</sup>MCS Division, Argonne National Laboratory, USA

<sup>2</sup>Institut für Mikroelektronik, TU Wien, Austria



C++Now, May 15th, 2013

Part 0: What is this all about?

# Contents

## Libraries

ViennaCL  
ViennaData  
ViennaFEM  
ViennaGrid  
ViennaIPD  
ViennaMath  
ViennaMesh

## Applications

ViennaMOS  
ViennaProfiler  
ViennaSHE  
ViennaWD  
ViennaX

# Contents

## Covered In This Talk

- Requirements in Computational Science

- GPU Computing

- Providing High-Level Interfaces

- Avoiding Monolithic Code

- Abstractions: From Math to Code (and back)

## Not Covered

- C++11

- Maximize use of Boost

# The Computational Scientist

## A Strange Animal

Goal is science, not to execute software

Codes seldomly designed for 'large scale' upfront

Scientists receive software training from scientists

Performance vs. portability and maintainability


Run code on clusters (headless, old software stack, ...)

*Come and go*

Basili *et al.*, Understanding the High-Performance-Computing Community: A Software Engineer's Perspective. IEEE Software, 2008.

# Simulation Flow

## The Many Steps in Simulating a FinFET




Gates

Fins

(C) Intel, 2011

# Simulation Flow


## The Many Steps in Simulating a FinFET



## Part 1: ViennaCL



# Simulation Flow



# From Boost.uBLAS to ViennaCL

## Consider Existing CPU Code (Boost.uBLAS)

```
using namespace boost::numeric::ublas;

matrix<double> A(1000, 1000);
vector<double> x(1000), y(1000);

/* Fill A, x, y here */

double val = inner_prod(x, y);
y += 2.0 * x;
A += val * outer_prod(x, y);

x = solve(A, y, upper_tag()); // Upper tri. solver

std::cout << " 2-norm: " << norm_2(x) << std::endl;
std::cout << "sup-norm: " << norm_inf(x) << std::endl
;
```

# From Boost.uBLAS to ViennaCL

## Previous Code Snippet Rewritten with ViennaCL

```
using namespace viennacl;
using namespace viennacl::linalg;

matrix<double> A(1000, 1000);
vector<double> x(1000), y(1000);

/* Fill A, x, y here */

double val = inner_prod(x, y);
y += 2.0 * x;
A += val * outer_prod(x, y);

x = solve(A, y, upper_tag()); // Upper tri. solver

std::cout << " 2-norm: " << norm_2(x) << std::endl;
std::cout << "sup-norm: " << norm_inf(x) << std::endl
;
```

# From Boost.uBLAS to ViennaCL

## ViennaCL in Addition Provides Iterative Solvers

```
using namespace viennacl;  
using namespace viennacl::linalg;  
  
compressed_matrix<double> A(1000, 1000);  
vector<double> x(1000), y(1000);  
  
/* Fill A, x, y here */  
  
x = solve(A, y, cg_tag());           // Conjugate  
    Gradients  
x = solve(A, y, bicgstab_tag()); // BiCGStab solver  
x = solve(A, y, gmres_tag());       // GMRES solver
```

## No Iterative Solvers Available in Boost.uBLAS...

# From Boost.uBLAS to ViennaCL

## Thanks to Interface Compatibility

```
using namespace boost::numeric::ublas;  
using namespace viennacl::linalg;  
  
compressed_matrix<double> A(1000, 1000);  
vector<double> x(1000), y(1000);  
  
/* Fill A, x, y here */  
  
x = solve(A, y, cg_tag());           // Conjugate  
    Gradients  
x = solve(A, y, bicgstab_tag()); // BiCGStab solver  
x = solve(A, y, gmres_tag());       // GMRES solver
```

## Code Reuse Beyond GPU Borders

Eigen <http://eigen.tuxfamily.org/>

MTL 4 <http://www.mtl4.org/>

# From Boost.uBLAS to ViennaCL

## Generic CG Implementation (Sketch)

```
for (unsigned int i = 0; i < tag.max_iterations(); ++
    i)
{
    tmp = viennacl::linalg::prod(A, p);

    alpha      = ip_rr / inner_prod(tmp, p);
    result     += alpha * p;
    residual   -= alpha * tmp;

    new_ip_rr = inner_prod(residual, residual);
    if (new_ip_rr / norm_rhs_squared < tag.tolerance())
        break;

    beta  = new_ip_rr / ip_rr;
    ip_rr = new_ip_rr;

    p = residual + beta * p;
}
```

# About ViennaCL


## About

High-level linear algebra C++ library

OpenMP, OpenCL, and CUDA backends

Header-only

Multi-platform



## Dissemination

Free Open-Source MIT (X11) License

<http://viennacl.sourceforge.net/>

50-100 downloads per week

## Design Rules

Reasonable default values

Compatible to Boost.uBLAS whenever possible

In doubt: clean design over performance

## Basic Types

scalar, vector

matrix, compressed\_matrix, coordinate\_matrix, ell\_matrix,  
hyb\_matrix

## Data Initialization

```
std::vector<double>          std_x(100);  
ublas::vector<double>       ublas_x(100);  
viennacl::vector<double>    vcl_x(100);  
  
for (size_t i=0; i<100; ++i)  
    // std_x[i] = rand();    // (1)  
    // ublas_x[i] = rand();  // (2)  
    vcl_x[i] = rand();      // (3)
```

(3) is slowest by orders of magnitude!



## Basic Types

scalar, vector

matrix, compressed\_matrix, coordinate\_matrix, ell\_matrix,  
hyb\_matrix

## Data Initialization

Using `viennacl::copy()`

```
std::vector<double>      std_x(100);  
ublas::vector<double>    ublas_x(100);  
viennacl::vector<double> vcl_x(100);  
  
/* setup of std_x and ublas_x omitted */  
  
viennacl::copy(std_x.begin(), std_x.end(),  
               vcl_x.begin()); //to GPU  
viennacl::copy(vcl_x.begin(), vcl_x.end(),  
               ublas_x.begin()); //to CPU
```

## Basic Types

scalar, vector

matrix, compressed\_matrix, coordinate\_matrix, ell\_matrix,  
hyb\_matrix

## Data Initialization

Using `viennacl::copy()`

```
std::vector<std::vector<double> >      std_A;  
ublas::matrix<double>                  ublas_A;  
viennacl::matrix<double>                vcl_A;  
  
/* setup of std_A and ublas_A omitted */  
  
viennacl::copy(std_A, vcl_A);           // CPU to GPU  
viennacl::copy(vcl_A, ublas_A);        // GPU to CPU
```

Iterator concept doesn't quite work on accelerators

## Vector Addition

```
x = y + z;
```

Temporaries are costly (particularly on GPUs)

## Expression Templates

Limited expansion

Map to a set of predefined kernels

```
vector_expression<vector<T>, op_plus, vector<T> >  
operator+(vector<T> & v, vector<T> & w) { ... }  
  
vector::operator=(vector_expression<...> const & e)  
{  
    viennacl::linalg::avbv(*this, 1,e.lhs(), 1,e.rhs()  
        );  
}
```


## Vector Addition

```
// x = y + z
void avbv(...) {
    switch (active_handle_id(x)) {
        case MAIN_MEMORY:
            host_based::avbv(...);
            break;
        case OPENCL_MEMORY:
            opencl::avbv(...);
            break;
        case CUDA_MEMORY:
            cuda::avbv(...);
            break;
        default:
            raise_error();
    }
}
```

Memory buffers can switch memory domain at runtime

## Memory Buffer Migration

```
vector<double> x = zero_vector<double>(42);  
  
memory_types src_memory_loc = memory_domain(x);  
switch_memory_domain(x, MAIN_MEMORY);  
  
/* work on x in main memory here */  
  
switch_memory_domain(x, src_memory_loc);
```



## Generalizing compute kernels

```
// x = y + z
__kernel void avbv(
    double * x,

    double * y,

    double * z, uint size)
{
    size_t i = get_global_id(0);
    for (; i < size; i += get_global_size())
        x[i] = y[i] + z[i];
}
```

## Generalizing compute kernels

```
// x = a * y + b * z
__kernel void avbv(
    double * x,
    double a,
    double * y,
    double b,
    double * z, uint size)
{
    size_t i = get_global_id(0);
    for (; i < size; i += get_global_size())
        x[i] = a * y[i] + b * z[i];
}
```

## Generalizing compute kernels

```
// x[4:8] = a * y[2:6] + b * z[3:7]
__kernel void avbv(
    double * x, uint off_x,
    double a,
    double * y, uint off_y,
    double b,
    double * z, uint off_z, uint size)
{
    size_t i = get_global_id(0);
    for (; i < size; i += get_global_size())
        x[off_x + i] = a * y[off_y + i] + b * z[off_z + i];
}
```




## Generalizing compute kernels


```
// x[4:2:8] = a * y[2:2:6] + b * z[3:2:7]
__kernel void avbv(
    double * x, uint off_x, uint inc_x,
    double a,
    double * y, uint off_y, uint inc_y,
    double b,
    double * z, uint off_z, uint inc_z, uint size)
{
    size_t i = get_global_id(0);
    for (; i < size; i += get_global_size())
        x[off_x + i * inc_x] = a * y[off_y + i * inc_y]
                               + b * z[off_z + i * inc_z];
}
```

No penalty on NVIDIA GPUs because FLOPs are for free

# Benchmarks




# Benchmarks



# Benchmarks

## Matrix-Matrix Multiplication

Autotuning environment



(AMD Radeon HD 7970, single precision)


# Expression Template Limitations

## Expression Templates are Not Enough

Consider

```
u = x + y;  
v = x - y;
```

Suboptimal performance with almost any library




# Expression Template Limitations

## Expression Templates are Not Enough

Consider

```
u = x + y;  
v = x - y;
```

Suboptimal performance with almost any library



# Expression Template Limitations

## Expression Templates are Not Enough

Consider

```
u = x + y;  
v = x - y;
```

Suboptimal performance with almost any library

## OpenCL Kernel Generation


Separate temporary avoidance from operation execution

```
custom_operation op;  
op.add( u = x + y );  
op.add( v = x - y );  
op.execute();           // OpenCL JIT
```

Fully transparent kernel fusion scheduled for release 1.5.0

# Expression Template Limitations


## Benchmark Results





## Part 2: ViennaGrid

# Simulation Flow



# What about Boost.Geometry?

## A Quick Look at Boost.Geometry

Concepts: Point, Segment, Linestring, Box, etc.

Algorithms: `distance()`, `intersects()`, `convex_hull()`, etc.

Central entity: Point

```
int a[3] = {1, 2, 3};  
int b[3] = {2, 3, 4};  
  
double d = boost::geometry::distance(a, b);
```

## Why Boost.Geometry is Not Enough

What about 3D objects? Tetrahedra? Hexahedra?


Traversal of boundary objects missing

No storage facilities for *many* objects

# $n$ -cell Concept

## Concept of $n$ -cell

Sub- $k$ -cells of an  $n$ -cell



## Separation of Geometry and Topology

Geometry: Euclidian space, coordinate system



Topology: Connection between points (lines, triangles, etc.)

# ViennaGrid Datastructure

## Datastructure Requirements

Don't store boundary  $k$ -cells unnecessarily

Fast local iteration on  $k$ -cells




	Amount	Mem/Obj.	Total Mem.
Vertices	4913	24 B	115 KB
Edges	31024	16 B	485 KB
Facets	50688	48 B	2376 KB
Cells	24576	112 B	2688 KB
<b>Total</b>			<b>5664 KB</b>

	Amount	Mem/Obj.	Total Mem.
Vertices	4913	24 B	115 KB
	0		0 KB

# *n*-cell Implementation

## Implementation of `element_t`

Recursive inheritance from boundary layer of dimension  $n - 1$   
Tag dispatching to enable/disable topological layers



```
template <typename ConfigType, typename ElementTag>
class element_t :
    public boundary_ncell_layer<ConfigType, ElementTag, ElementTag::
        dim-1>,
    public result_of::element_id_handler<ConfigType, ElementTag>::
        type
```

# Domain Configuration

## Top Level Configuration of Domain

Mostly a collection of tags

Predefined configuration classes for common cases

ViennaGrid 1.0.x: Supplemented by global customizations

```
struct triangular_3d
{
    typedef double          numeric_type;
    typedef cartesian_cs<3> coordinate_system_tag;
    typedef triangle_tag    cell_tag;
};

result_of::domain<triangular_3d>::type    Domain;
```

Work in progress: Everything in a single config class

## User API Design Goals

STL-style, reuse conventions

Allow index-based traversal

Avoid common C++ pitfalls (e.g. template member functions)

```
//iteration over all vertices in the domain:
typedef result_of::ncell_range<DomainType, 0>::type   VertexRange;
typedef result_of::iterator<VertexRange>::type       VertexIterator
    ;

VertexRange vertices = ncells<0>(domain);
for (VertexIterator it  = vertices.begin();
      it != vertices.end();
      ++it)
{
    // do something with each vertex here
}
```

*Ranges* provide iterators over  $n$ -cells

Extendible



## User API Design Goals

STL-style, reuse conventions

Allow index-based traversal

Avoid common C++ pitfalls (e.g. template member functions)

```
//iteration over all vertices in the domain:  
for (std::size_t i = 0; i < ncells<0>(domain).size(); ++i)  
{  
    // do something with ncells<0>(domain)[i] here  
}
```


OpenMP friendly

.size() sometimes known at compiletime!

## User API Design Goals

Boundary iteration:  $k < n$

Coboundary iteration:  $k > n$



```
//iteration over all triangles of a vertex
typedef result_of::ncell_range<VertexType, 2>::type TriangleRange;
typedef result_of::iterator<TriangleRange>::type
    TriangleIterator;

TriangleRange triangles = ncells<2>(vertex, domain);
for (TriangleIterator it = triangles.begin();
      it != triangles.end();
      ++it)
{
    // do something with each triangle here
}
```

Coboundary information not a-priori available in datastructure  
Built and cached at first request

# ViennaGrid Features

## Other Features

- Segments


- I/O: VTK, various mesh-format

- Voronoi information

- Refinement


## Work in Progress

- PLC, hybrid meshes, multigrid



## Part 3: ViennaData

# Simulation Flow



## Plain Object-Oriented Approach?

```
struct Triangle {  
    PointType a, b, c;  
    bool is_on_boundary;  
    double rho; };    //e.g. specific mass
```

## Pros and Cons

Data is directly stored with the object

Each (!) triangle carries a boolean flag

Reusability reduced  $\Rightarrow$  better rename to

TriangleWithMaterial

# Monolithic, don't do this!

# Data Storage

## Store Data with Mesh:

```
class Mesh {  
    vector<Triangle> triangles;  
    std::map<size_t, bool> boudary_map; //no memory  
    wasted  
    vector<double> rho_for_triangles;    };
```

## Pros and Cons

Triangle is still reusable

Mesh class has to handle data storage complexity

Each additional data requires a change of Mesh class

Mesh object has to be passed to all modules

**Monolithic, don't do this!**

## Approach by ViennaData

Introduction of a hidden data container

Data is stored in a map-like manner using keys

```
//generic interface:  
viennadata::access<KeyType, ValueType>(key) (obj);  
  
//boundary flag and specific mass retrieval for triangle:  
bool on_boundary = access<BoundaryKeyType, bool>(boundary_key) (triangle);  
double rho = access<MassKeyType, double>( mass_key) (triangle);
```

## Pros and Cons

- + Triangle and Mesh are still reusable
- + Can be used with arbitrary objects (third-party libraries)
- + Unified interface for data access
- Global state



## Generic Interface

```
viennadata::access<KeyType, ValueType>(key) (obj);
```

## Default Storage Scheme

If nothing is known about the object:

```
std::map<ObjectType *,  
        std::map<KeyType, ValueType> >
```

The compiler creates such a map for each of the following:

```
access<long, double>(42) (triangle);  
access<char, double>('b') (triangle);  
access<std::string, double>("mass") (triangle);  
access<std::string, double>("mass") (vertex);
```

## Performance Considerations

$\mathcal{O}(\log N) + \mathcal{O}(\log K)$  access time

$N$  ... objects of same type (e.g. triangles)

$K$  ... keys of same type

Usually too slow in a high-performance setting

## Can We Do Better?

In general: No

In certain situations: Yes

## Type-based Key Dispatch

Only one key per type

$\mathcal{O}(\log N)$  access time

```
access<mass_key, double>(mass_key())(triangle);
```

```
// or:
```

```
access<mass_key, double>()(triangle);
```

## Internal Datastructure

```
std::map<ObjectType *, ValueType>
```

Transparent to user

One line of code for activation

## Numeric IDs for Objects

Only one key per type

$\mathcal{O}(K)$  access time

$\mathcal{O}(1)$  access time with type-based key dispatch

```
access<std::string, double>("mass") (triangle);  
  
// with type-based key dispatch:  
access<mass_key, double>() (triangle);
```

## Internal Datastructure

```
std::vector<ValueType>
```

Transparent to user

One line of code for activation

Overload generic `id()` accessor

## Benchmarking ViennaData

ID-based access to data via ViennaData (class `LightWeight`)  
OOP-style storage in classes with payload

```
template <size_t N>
struct FatClass {
    double data;
    char payload[N];
};
```

	10 <sup>3</sup> Objects (us)	10 <sup>6</sup> Objects (ms)
LightWeight	<b>4</b>	<b>5</b>
FatClass<10>	<b>1.3</b>	<b>4</b>
FatClass<100>	2.1	11
FatClass<1000>	2.5	11

## Other Routines

Not further addressed:

```
viennadata::copy<KeyType, ValueType>(key) (  
    src_obj, dst_obj);  
viennadata::move<KeyType, ValueType>(key) (  
    src_obj, dst_obj);  
viennadata::find<KeyType, ValueType>(key) (obj);  
  
viennadata::erase<KeyType, ValueType>(key) (obj);  
viennadata::erase<KeyType, all>(key) (obj);  
viennadata::erase<all, ValueType>(key) (obj);
```

## Pitfalls

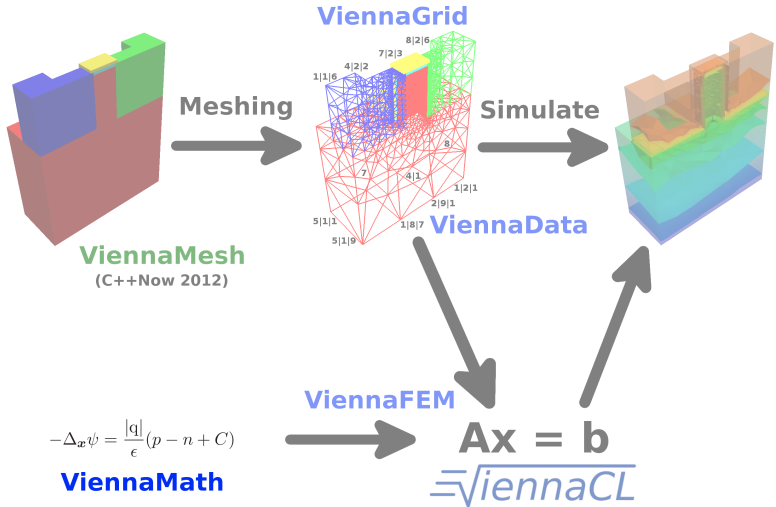
Inheritance

Limited lifetime of objects

Compilation units with different configuration

## Part 4: ViennaMath

# Simulation Flow





## A Symbolic Math Library in C++

Symbolic evaluation and manipulation of math expressions

Unified run time and compile time interface

Differentiation and integration provided

$\text{\LaTeX}$  output

## Example Usage

```
variable x(0);  
variable y(1);  
variable z(2);  
expr f = x + y - z;  
expr g = f * f;  
eval(f, make_vector(1.0, 2.0, 4.0)); // returns  
-1.0  
eval(g, make_vector(1.0, 2.0, 4.0)); // returns  
1.0
```

## A Symbolic Math Library in C++

Symbolic evaluation and manipulation of math expressions

Unified run time and compile time interface

Differentiation and integration provided

$\text{\LaTeX}$  output

## Example Usage

```
ct_variable<0> x;  
ct_variable<1> y;  
ct_variable<2> z;  
ct_constant<1> c1;  
ct_constant<2> c2;  
ct_constant<4> c4;  
eval(x + y - z, make_vector(c1, c2, c4)); //  
    returns -1
```

## Substitution and Differentiation (Run Time)

```
variable x(0);  
variable y(1);  
variable z(2);  
substitute(x, y, x*y + z); // returns y*y+z  
diff(x*y + z, x);          // returns y
```

## Substitution and Differentiation (Compile Time)

```
ct_variable<0> x;  
ct_variable<1> y;  
ct_variable<2> z;  
substitute(x, y, x*y + z); // returns y*y+z  
diff(x*y + z, x);          // returns y
```

## Numerical Integration (Run Time): $\int_0^1 x^2 dx$

```

expr f = integral( make_interval(0, 1), x*x, x );

numerical_quadrature integrator(new gauss_quad_1())
;
integrator(f);                                //
    method 1
integrator(make_interval(0, 1), x*x, x);      //
    method 2

```

## Analytical Integration (Compile Time): $\int_0^1 x^2 dx$ , $\int_0^1 \int_0^{1-x} xy dx dy$

```

integrate(make_interval(c0, c1),
          x*y,
          x );    //returns y/2.0
integrate(make_interval(c0, c1),
          integrate( make_interval(c0, c1 - x), x*y

```

## Function Symbols

Represent a function (not evaluable)

## Differential Operators

Gradient, Divergence


## Symbolic Integration Domain


Specify actual integration domain and integration variables  
*later*

```
function_symbol u;  
equation eq( laplace(u), 1.0 ); // Poisson  
equation  
  
function_symbol v;  
expr w = integral(symbolic_interval(),  
                   grad(u) * grad(v));
```

## Part 5: ViennaFEM

# Simulation Flow





## Library-Centric Design

ViennaCL for linear solver

ViennaData for data storage

ViennaGrid for mesh handling

ViennaMath for symbolic math

## Addresses Come&Go in Academia

Focus on one package

No in-depth knowledge of all packages required

Additional emphasis on good interfaces



## ViennaGrid deals with the Mesh

```
DomainType my_domain;  
viennagrid::io::netgen_reader my_reader;  
my_reader(my_domain, "mystructure.mesh");
```

## Equation specification via ViennaMath: $\Delta u = -1$

```
equation poisson_eq = viennamath::make_equation(laplace(u), -1);
```

## Assembly via ViennaFEM

```
viennafem::pde_assembler fem_assembler;  
fem_assembler(viennafem::make_linear_pde_system(poisson_eq, u),  
              my_domain,  
              system_matrix, load_vector);
```

## Linear solver provided by ViennaCL

```
VectorType pde_result  
= viennacl::linalg::solve(system_matrix, load_vector, cg_tag());
```

## Lame equation for Linear Elasticity


$$\int_{\Omega} \varepsilon(u) : \sigma(v) \, d\Omega = \int_{\Omega} F \cdot v \, d\Omega \quad \forall v \in \mathcal{V}$$

With  $F = (0, 0, 1)^T$ :

```
std::vector< Expression > strain = strain_tensor(u)
;
std::vector< Expression > stress = stress_tensor(v)
;

Equation weak_form_lame = make_equation(
    integral(symbolic_interval(),
              tensor_reduce( strain, stress )
            ),
    // =
    integral(symbolic_interval(),
              constant(1.0) * v[2])
)
```

# Solving Lamé's Equation for a TSV



# Summary

## Software Packages

- ViennaCL: GPU-accelerated Linear Algebra
- ViennaData: Generic Data Storage
- ViennaFEM: Modular High-Level Finite Element Package
- ViennaGrid: Generic Mesh Datastructure
- ViennaMath: Symbolic Math Kernel

`http://vienna{cl,data,fem,grid,math}@sourceforge.net`

## Design Philosophy

- Orthogonal Software Design
- Convenient High-Level User API
- Avoid Unneeded Dependencies
- Free Open Source (MIT License)