

Exascale Computing Without Templates

Karl Rupp, Richard Mills, Barry Smith, Matthew Knepley, Jed Brown



Argonne National Laboratory



DOE COE Performance Portability Meeting, Denver
August 22-24, 2017

Exascale Computing without Threads*

A White Paper Submitted to the
DOE High Performance Computing Operational Review (HPCOR)
on Scientific Software Architecture for Portability and Performance
August 2015

Matthew G. Knepley¹, Jed Brown², Barry Smith², Karl Rupp³, and Mark Adams⁴

¹Rice University, ²Argonne National Laboratory, ³TU Wien, ⁴Lawrence Berkeley National Laboratory
knepley@rice.edu, [jedbrown,bsmith]@mcs.anl.gov, rupp@iue.tuwien.ac.at, mfadams@lbl.gov

Abstract

We provide technical details on why we feel the use of threads does not offer any fundamental performance advantage over using processes for high-performance computing and hence why we plan to extend PETSc to exascale (on emerging architectures) using node-aware MPI techniques, including neighborhood collectives and portable shared memory within a node, instead of threads.

https://www.ornl.gov/hpcor2015/whitepapers/Exascale_Computing_without_Threads-Barry_Smith.pdf

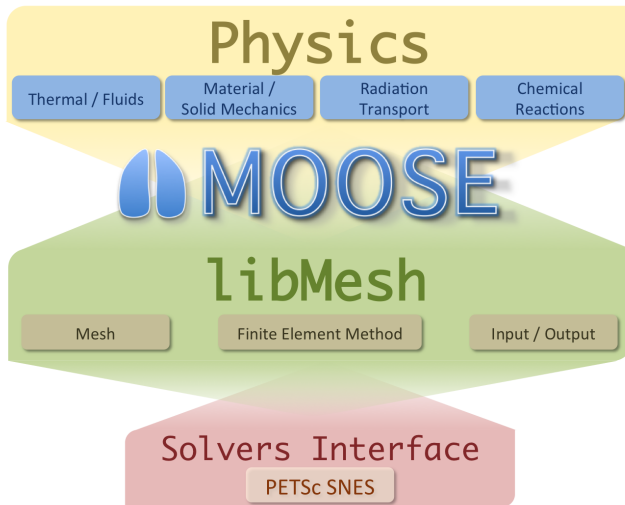
PETSc Developers Care About Recent Developments

After careful evaluation: Favor MPI 3.0 (and later) over threads

Find the best long-term solutions for our users

Consider best solutions for large-scale applications, not just toy-apps

Providing Context



Our Attempts in C++ Library Development

Sieve: Several years of C++ mesh management attempts in PETSc

ViennaGrid 2.x: Heavily templated C++ mesh management library

ViennaCL: Dense and sparse linear algebra and solvers for multi- and many-core architectures


Aftermath

Sieve: Replaced by DMPlex (written in C)

ViennaGrid: Version 3.0 provides C-ABI

ViennaCL: Rewrite in C likely

Sequential build times for the ViennaCL test suite



Disadvantages of C++ Templates

Static Dispatch

Architecture-specific information only available at run time

“Change code and recompile” not acceptable advice


Dealing with Compilation Errors

Type names pollute compiler output

Replicated across interfaces

CRTP may result in type length explosion

Default arguments become visible



<https://xkcd.com/303/>

Type	Length
<code>std::vector<int></code>	38
<code>std::vector<std::vector<int> ></code>	109
<code>std::vector<std::vector<std::vector<int> > ></code>	251
<code>std::vector<std::vector<std::vector<std::vector<int> > > ></code>	539



<https://tgceec.tumblr.com/>

Scope Limitations

Template metaprogramming lacks state

Optimizations across multiple code lines difficult or impossible

Example

Consider vector updates in pipelined CG method:

$$x_i \leftarrow x_{i-1} + \alpha p_{i-1}$$

$$r_i \leftarrow r_{i-1} - \alpha y_i$$

$$p_i \leftarrow r_i + \beta p_{i-1}$$

Reuse of p_{i-1} and r_{i-1} easy with for-loops, but hard with expression templates

Complicates Debugging

- Stack traces get longer names and deeper
- Setting good breakpoints may become harder

Lack of a Stable ABI

- Object files from different compilers generally incompatible
- Name mangling makes use outside C++ land almost impossible

High Entry Bar

- Number of potential contributors inversely proportional to code sophistication
- Domain scientists have limited resources for C++ templates

Manage Complexity

Good interface design

Refactor code when needed

Hand-optimize small kernels only (cf. BLIS methodology)

Development Implications


Adopt professional software development practices

Develop, maintain, and evolve different datastructures ...

... and code paths

Use clear and easy-to-understand datastructures

Fallacy: “Writing” an application only once in its final form



Spending Development Resources

- Reuse existing libraries — reinventing the wheel is not productive!

- Focus on domain- and application-specific aspects

- Obtain expertise and resources for continuous code evolution


Required Incentives

- Reward contributions to existing projects

- Pair research funding with software development funding

- Establish software development career tracks

Is Performance Portability Just a Software Productivity Aspect?



<https://www.nitrd.gov/PUBS/CSESSPWorkshopReport.pdf>

Long-Term Problems of Heavy C++ Templates Use

Template metaprogramming is a leaky abstraction

Excessive type names slow down all stages of Compile-Run-Debug-cycle

Templates operate at compile time - architecture ultimately known at run time

A Path Forward

Adopt professional software development practices

Be prepared to develop different datastructures and code paths

Write clear, readable code using simple datastructures

Evolve and refactor datastructures, kernels, and interfaces over time

(cf. software productivity discussions)