

Parallel Solver Essentials for Computational Scientists

Part 1: Shared Memory Systems

Karl Rupp

`me@karlrupp.net`

`https://karlrupp.net/`

`http://github.com/karlrupp/slides`

 `@karlrupp`

now:

Freelance Scientist

formerly and soon again:

Institute for Microelectronics, TU Wien



Current Many-Core Architectures


High FLOP/Watt ratio

High memory bandwidth

Attached via PCI-Express



AMD FirePro W9100
320 GB/sec




INTEL Xeon Phi
320 (220?) GB/sec



NVIDIA Tesla K20
250 (208) GB/sec

Introduction




Original data up to the year 2010 collected and plotted by M. Horowitz, F. Labonte, O. Shacham, K. Olukotun, L. Hammond, and C. Batten
New plot and data collected for 2010-2015 by K. Rupp

Introduction

Theoretical Peak Performance

Theoretical Peak Performance, Double Precision

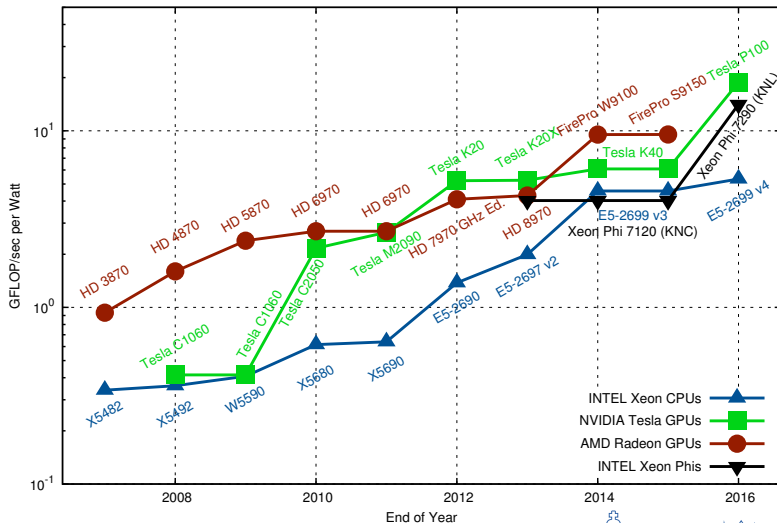


<https://www.karlrupp.net/2013/06/cpu-gpu-and-mic-hardware-characteristics-over-time/>

Introduction


Theoretical Peak Performance per Watt

Theoretical Peak Floating Point Operations per Watt, Double Precision



Introduction

Theoretical Peak Performance (FLOPs) per Byte of Memory Bandwidth



About

Initial release in 2007

Proprietary programming model by NVIDIA

C++ with extensions

Proprietary compiler extracts GPU kernels

Software Ecosystem

Vendor-tuned libraries: cuBLAS, cuSparse, cuSolver, cuFFT, etc.

Python bindings: pyCUDA

Community projects: CUSP, MAGMA, VexCL, ViennaCL, etc.



Programming in CUDA

```
void work(double *x, double *y, double *z, int N)
{

    for (size_t i=0; i<N; ++i)
        z[i] = x[i] + y[i];
}
```

```
int main(int argc, char **argv)
{
    int N = atoi(argv[1]);
    double *x = malloc(N*sizeof(double));
    ...

    ...
    work(x, y, z, N); // call kernel
    ...

    free(x);
}
```


Programming in CUDA

```
void work(double *x, double *y, double *z, int N)
{
    #pragma omp parallel for
    for (size_t i=0; i<N; ++i)
        z[i] = x[i] + y[i];
}
```

```
int main(int argc, char **argv)
{
    int N = atoi(argv[1]);
    double *x = malloc(N*sizeof(double));
    ...

    ...
    work(x, y, z, N); // call kernel
    ...

    free(x);
}
```

Programming in CUDA

```
void work(double *x, double *y, double *z, int N)
{
    #pragma omp parallel
    { int thread_id = omp_get_thread_num();
      for (size_t i=thread_id; i<N; i += omp_get_num_threads())
          z[i] = x[i] + y[i];
    } }
```

```
int main(int argc, char **argv)
{
    int N = atoi(argv[1]);
    double *x = malloc(N*sizeof(double));
    ...



    ...
    work(x, y, z, N); // call kernel
    ...

    free(x);
}
```

Programming in CUDA

```
__global__ void work(double *x, double *y, double *z, int N)
{
    int thread_id = blockIdx.x*blockDim.x + threadIdx.x;
    for (size_t i=thread_id; i<N; i += blockDim.x * gridDim.x)
        z[i] = x[i] + y[i];
}
```

```
int main(int argc, char **argv)
{
    int N = atoi(argv[1]);
    double *x = malloc(N*sizeof(double));
    cudaMalloc(&gpu_x, N*sizeof(double));
    cudaMemcpy(gpu_x, x, N*8, cudaMemcpyHostToDevice);
    ...
    work<<<128, 256>>>(x, y, z, N); // call kernel
    ...
    cudaMemcpy(gpu_x, x, N*8, cudaMemcpyDeviceToHost);
    ...
    free(x);
}
```



Thread Control (1D)

Local ID in block: `threadIdx.x`

Threads per block: `blockDim.x`

ID of block: `blockIdx.x`

No. of blocks: `gridDim.x`

Recommended Default Values

Typical block size: 256 or 512

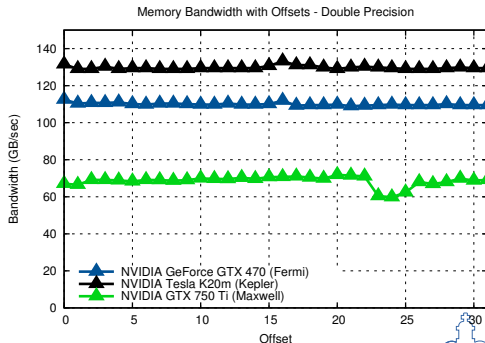
Typical number of blocks: 256

At least 10 000 logical threads recommended



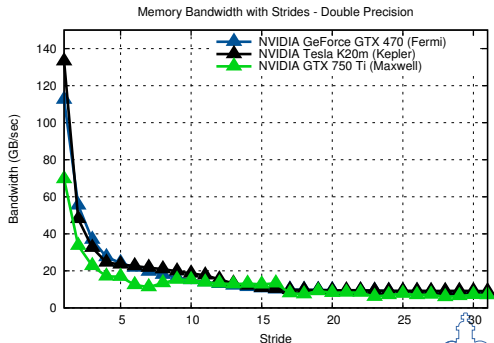
Offset Memory Access

```
__global__  
void work(double *x, double *y, double *z, int N, int k)  
{  
    int thread_id = blockIdx.x*blockDim.x + threadIdx.x;  
    for (size_t i=thread_id; i<N; i += blockDim.x * gridDim.x)  
        z[i+k] = x[i+k] + y[i+k];  
}
```



Strided Memory Access

```
__global__  
void work(double *x, double *y, double *z, int N, int k)  
{  
    int thread_id = blockIdx.x*blockDim.x + threadIdx.x;  
    for (size_t i=thread_id; i<N; i += blockDim.x * gridDim.x)  
        z[i*k] = x[i*k] + y[i*k];  
}
```




Strided Memory Access

Array of structs problematic

```
typedef struct particle
{
    double pos_x; double pos_y; double pos_z;
    double vel_x; double vel_y; double vel_z;
    double mass;
} Particle;

__global__
void increase_mass(Particle *particles, int N)
{
    int thread_id = blockIdx.x*blockDim.x + threadIdx.x;
    for (int i=thread_id; i<N; i += blockDim.x * gridDim.x)
        particles[i].mass *= 2.0;
}
```




Strided Memory Access

Workaround: Structure of Arrays

```
typedef struct particles
{
    double *pos_x; double *pos_y; double *pos_z;
    double *vel_x; double *vel_y; double *vel_z;
    double *mass;
} Particle;

__global__
void increase_mass(Particle *particles, int N)
{
    int thread_id = blockIdx.x*blockDim.x + threadIdx.x;
    for (int i=thread_id; i<N; i += blockDim.x * gridDim.x)
        particles.mass[i] *= 2.0;
}
```



Reductions

Use N values to compute 1 result value


Examples: Dot-products, vector norms, etc.

Reductions with Few Threads

Decompose N into chunks for each thread

Compute chunks in parallel

Merge results with single thread




Reductions with Many Threads


Decompose N into chunks for each workgroup

Use fast on-chip synchronization within each workgroup

Sum result for each workgroup separately



Reductions with Many Threads



```
shared_m[threadIdx.x] = thread_sum;
for (int stride = blockDim.x/2; stride>0; stride/=2) {
    __syncthreads();
    if (threadIdx.x < stride)
        shared_m[threadIdx.x] += shared_m[threadIdx.x+stride];
}
```

Prefix Sum

Inclusive: Determine $y_i = \sum_{k=1}^i x_k$

Exclusive: Determine $y_i = \sum_{k=1}^{i-1} x_k, y_1 = 0$

Example

x: 4, 3, 6, 5, 4, 7, 4, 4, 4


y: 4, 7, 13, 18, 22, 29, 33, 37, 41 (inclusive)

y: 0, 4, 7, 13, 18, 22, 29, 33, 37 (exclusive)


Applications

Sparse matrix setup

Graph algorithms



Prefix Sum Implementation



```
for(int stride = 1; stride < blockDim.x; stride *= 2)
{
    __syncthreads();
    shared_buffer[threadIdx.x] = my_value;
    __syncthreads();
    if (threadIdx.x >= stride)
        my_value += shared_buffer[threadIdx.x - stride];
}
__syncthreads();
shared_buffer[threadIdx.x] = my_value;
```

Other Parallel Primitives

Sort

Gather and Scatter

Load to shared memory and work there

etc.

GPU-Accelerated Software Libraries

Linear Algebra: ViennaCL, MAGMA, CUSP, VexCL, ...

Solvers: ViennaCL, MAGMA, cuSolver, Paralution, clAMG, ...

FFT: cuFFT, clFFT, FFTW, ...

Primitives: VexCL, Boost.Compute, ...

Machine Learning: Caffe, cuDNN, ...



Pipelined CG

- Merge global reductions
- Kernel fusion

Parallel Incomplete LU Factorizations

- Level scheduling
- Nonlinear relaxation

Algebraic Multigrid

- Parallel aggregation
- Sparse matrix-matrix products



Pseudocode

Choose x_0

$$p_0 = r_0 = b - Ax_0$$

For $i = 0$ until convergence

1. Compute and store Ap_i
2. Compute $\langle p_i, Ap_i \rangle$
3. $\alpha_i = \langle r_i, r_i \rangle / \langle p_i, Ap_i \rangle$
4. $x_{i+1} = x_i + \alpha_i p_i$
5. $r_{i+1} = r_i - \alpha_i Ap_i$
6. Compute $\langle r_{i+1}, r_{i+1} \rangle$
7. $\beta_i = \langle r_{i+1}, r_{i+1} \rangle / \langle r_i, r_i \rangle$
8. $p_{i+1} = r_{i+1} + \beta_i p_i$

EndFor

BLAS-based Implementation


-

SpMV, AXPY


For $i = 0$ until convergence

1. SpMV \leftarrow No caching of Ap_i
2. DOT \leftarrow Global sync!
3. -
4. AXPY
5. AXPY \leftarrow No caching of r_{i+1}
6. DOT \leftarrow Global sync!
7. -
8. AXPY

EndFor



Performance Modeling: Conjugate Gradients



(Poisson, 2D, Finite Differences)


Performance Modelling

6 Kernel Launches (plus two for reductions)

Two device to host data reads from dot products

Model SpMV as seven vector accesses (5-point stencil)

$$T(N) = 8 \times 10^{-6} + 2 \times 2 \times 10^{-6} + (7 + 2 + 3 + 3 + 2 + 3) \times 8 \times N / \text{Bandwidth}$$



Optimization: Rearrange the algorithm

- Remove unnecessary reads

- Remove unnecessary synchronizations

- Use custom kernels instead of standard BLAS



Standard CG

Choose x_0

$$p_0 = r_0 = b - Ax_0$$

For $i = 0$ until convergence

1. Compute and store Ap_i
2. Compute $\langle p_i, Ap_i \rangle$
3. $\alpha_i = \langle r_i, r_i \rangle / \langle p_i, Ap_i \rangle$
4. $x_{i+1} = x_i + \alpha_i p_i$
5. $r_{i+1} = r_i - \alpha_i Ap_i$
6. Compute $\langle r_{i+1}, r_{i+1} \rangle$
7. $\beta_i = \langle r_{i+1}, r_{i+1} \rangle / \langle r_i, r_i \rangle$
8. $p_{i+1} = r_{i+1} + \beta_i p_i$

EndFor

Pipelined CG


Choose x_0

$$p_0 = r_0 = b - Ax_0$$


For $i = 1$ until convergence

1. $i = 1$: Compute α_0, β_0, Ap_0
2. $x_i = x_{i-1} + \alpha_{i-1} p_{i-1}$
3. $r_i = r_{i-1} - \alpha_{i-1} Ap_{i-1}$
4. $p_i = r_i + \beta_{i-1} p_{i-1}$
5. Compute and store Ap_i
6. Compute $\langle Ap_i, Ap_i \rangle, \langle p_i, Ap_i \rangle, \langle r_i, r_i \rangle$
7. $\alpha_i = \langle r_i, r_i \rangle / \langle p_i, Ap_i \rangle$
8. $\beta_i = (\alpha_i^2 \langle Ap_i, Ap_i \rangle - \langle r_i, r_i \rangle) / \langle r_i, r_i \rangle$

EndFor




Performance Modeling: Conjugate Gradients



(Poisson, 2D, Finite Differences)

Benefits of Pipelining also for Large Matrices



Parallel Incomplete LU Factorizations

Level scheduling

Nonlinear relaxation



ILU - Basic Idea

Factor sparse matrix $A \approx \tilde{L}\tilde{U}$

\tilde{L} and \tilde{U} sparse, triangular

ILU0: Pattern of \tilde{L} , \tilde{U} equal to A

ILUT: Keep k elements per row

Solver Cycle Phase

Residual correction $\tilde{L}\tilde{U}x = z$

Forward solve $\tilde{L}y = z$

Backward solve $\tilde{U}x = y$

Little parallelism in general

$$\begin{pmatrix} 5 & \times & \times & \times & & \times & \times & & \\ \times & 3 & \times & & & & & & \\ \times & \times & 4 & \times & & & & & \\ \times & & \times & 5 & \times & \times & & & \times \\ & & & \times & 5 & \times & & \times & \times \\ \times & & & \times & \times & 6 & \times & \times & \\ \times & & & & & \times & 3 & & \\ & & & & \times & \times & & 4 & \times \\ & & \times & \times & & & & \times & 4 \end{pmatrix}$$



ILU Level Scheduling

Build dependency graph

Substitute as many entries as possible simultaneously

Trade-off: Each step vs. multiple steps in a single kernel



$$\begin{pmatrix}
 5 & \times & \times & \times & & \times & \times & & \\
 \times & 3 & \times & & & & & & \\
 \times & \times & 4 & \times & & & & & \\
 \times & \times & \times & 5 & \times & \times & & & \times \\
 \times & & \times & \times & 5 & \times & & \times & \times \\
 \times & & \times & \times & \times & 6 & \times & \times & \\
 \times & & \times & \times & \times & \times & 3 & & \\
 \times & & \times & \times & \times & \times & \times & 4 & \times \\
 \times & & \times & \times & \times & \times & \times & \times & 4
 \end{pmatrix}$$

ILU Interpretation on Structured Grids

2d finite-difference discretization

Substitution whenever all neighbors with smaller index computed

Works particularly well in 3d





ILU Interpretation on Structured Grids

2d finite-difference discretization

Substitution whenever all neighbors with smaller index computed

Works particularly well in 3d





ILU Interpretation on Structured Grids

2d finite-difference discretization

Substitution whenever all neighbors with smaller index computed

Works particularly well in 3d





ILU Interpretation on Structured Grids

2d finite-difference discretization

Substitution whenever all neighbors with smaller index computed

Works particularly well in 3d



Sequential

```
for i=2..n
  for k=1..i-1, (i,k) in A
     $a_{ik} = a_{ik}/a_{kk}$ 
  for j=k+1..n, (i,j) in A
     $u_{ij} = a_{ij} - a_{ik}a_{kj}$ 
```

Parallel

```
for (sweep = 1, 2, ...)
  parallel for (i,j) in A
    if (i > j)
       $l_{ij} = (a_{ij} - \sum_{k=1}^{j-1} l_{ik}u_{kj})/u_{jj}$ 
    else
       $u_{ij} = a_{ij} - \sum_{k=1}^{i-1} l_{ik}u_{kj}$ 
```

Fine-Grained Parallel ILU Setup


Proposed by Chow and Patel (SISC, vol. 37(2)) for CPUs and MICs
Massively parallel (one thread per row)

Preconditioner Application


Truncated Neumann series:

$$\mathbf{L}^{-1} \approx \sum_{k=0}^K (\mathbf{I} - \mathbf{L})^k, \quad \mathbf{U}^{-1} \approx \sum_{k=0}^K (\mathbf{I} - \mathbf{U})^k$$

Exact triangular solves not necessary



Parallel ILU



Algebraic Multigrid

Parallel aggregation

Sparse matrix-matrix products




Ingredients of Algebraic Multigrid


Smoother (Relaxation schemes, etc.)

Coarsening

Interpolation (Inter-grid transfer)



Classical coarsening



Aggregation coarsening

Setup Phase

Determination of coarse points in parallel by graph splitting

Compute coarse operators $A^{k+1} = R^k A^k P^k$ (where $A^0 = A$)

Datastructures: analyze and allocate

Limited fine-grained parallelism

Cycle Phase


Parallel Jacobi Smoother

Restriction $R^k \chi^k$, prolongation $P^k \chi^{k+1}$

Direct solution on coarsest level

Static datastructures

Enough fine-grained parallelism



Coarse Grid Operator

$$A^{\text{coarse}} = RA^{\text{fine}}P$$

Common choice: $R = P^T$

Computation


Explicitly set up $R = P^T$ (hard in parallel)

$$C = A^{\text{fine}}P$$

$$A^{\text{coarse}} = RC$$




AMG Sparse Matrix-Matrix Multiplication



AMG Benchmark

Total Solver Execution Times, Poisson Equation in 2D



Parallel Primitives

- Embarassingly parallel operations (vector addition, etc.)

- Reductions

- Prefix Sums

- etc.

Solvers on Shared Memory Architectures

- Pipelining to reduce synchronization costs

- Pipelining to increase data reuse

- Replace sequential stages with parallel alternatives

