

Performance Tuning for GPUs - An Iterative Process

Karl Rupp^{1,2}

`rupp@iue.tuwien.ac.at`

 @karlrupp

with contributions from

Philippe Tillet¹, Florian Rudolf¹,
Josef Weinbub¹, Ansgar Jüngel², Tibor Grasser¹
(based on stimuli from PETSc+ViennaCL users)



¹ Institute for Microelectronics, TU Wien, Austria

² Institute for Analysis and Scientific Computing, TU Wien, Austria


Introduction



Conjugate Gradient - Standard Formulation



Conjugate Gradient - Kernel Parameters



Parameter Study for Portable Performance



Conjugate Gradient - Pipelining



GMRES Optimizations



Conclusion



Positions

PhD student at TU Wien (2009-2011)

Postdoc at ANL (09/2012-09/2013)

Postdoc at TU Wien (01/2012-09/2012, 09/2013-current)

Research Interests

Semiconductor device simulation

Numerical solution of PDEs

Parallel computing

Software Development

PETSc

ViennaCL

ViennaSHE

...



Iterative Solvers

Matrix-vector products and vector operations only


Expose more fine-grained parallelism

Preconditioners often desirable


Accelerators (CUDA, OpenCL)

Graphics processing units (GPUs)

Intel Xeon Phi



Introduction



Pseudocode

Choose x_0

$$p_0 = r_0 = b - Ax_0$$

For $i = 0$ until convergence

1. Compute and store Ap_i
2. Compute $\langle p_i, Ap_i \rangle$
3. $\alpha_i = \langle r_i, r_i \rangle / \langle p_i, Ap_i \rangle$
4. $x_{i+1} = x_i + \alpha_i p_i$
5. $r_{i+1} = r_i - \alpha_i Ap_i$
6. Compute $\langle r_{i+1}, r_{i+1} \rangle$
7. $\beta_i = \langle r_{i+1}, r_{i+1} \rangle / \langle r_i, r_i \rangle$
8. $p_{i+1} = r_{i+1} + \beta_i p_i$

EndFor

BLAS-based Implementation

-

SpMV, AXPY


For $i = 0$ until convergence

1. SpMV \leftarrow No caching of Ap_i
2. DOT \leftarrow Global sync!
3. -
4. AXPY
5. AXPY \leftarrow No caching of r_{i+1}
6. DOT \leftarrow Global sync!
7. -
8. AXPY

EndFor



Conjugate Gradients





(2D Finite Difference Discretization)

Implications

Kernel launches expensive

Delicate balance for preconditioners

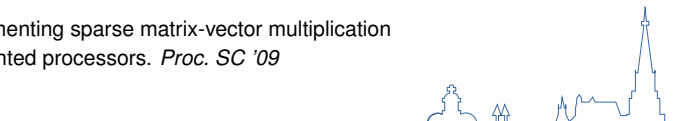


Optimization 1


Get best performance out of SpMV

Compare different sparse matrix types

Cf.: N. Bell: Implementing sparse matrix-vector multiplication on throughput-oriented processors. *Proc. SC '09*




Conjugate Gradients



(2D Finite Difference Discretization)


Conjugate Gradients



(2D Finite Difference Discretization)

Optimization 2

Optimize kernel parameters for each operation




Introduction



Conjugate Gradient - Standard Formulation



Conjugate Gradient - Kernel Parameters



Parameter Study for Portable Performance



Conjugate Gradient - Pipelining



GMRES Optimizations




Conclusion



Scope for OpenCL-based Portability Study

Vector and matrix-vector operations (BLAS levels 1 and 2)

Limited by memory bandwidth




Scope for OpenCL-based Portability Study

Vector and matrix-vector operations (BLAS levels 1 and 2)

Limited by memory bandwidth


Key Question (Memory-Bandwidth-Limited Kernels)

Good performance of complicated kernels
by optimizing the simplest kernel?



Vector Assignment (Copy) Kernel

$x \leftarrow y$ for (large) vectors x, y



Parameters (1900 variations)

Local work size, global work size


Vector types (float1, float2, ... , float16)

Thread increment type




Vector Assignment (Copy) Kernel

$x \leftarrow y$ for (large) vectors x, y




Parameters (1900 variations)

```
for (size_t i = get_global_id(0); i < N;  
      i += get_global_size(0))  
    x[i] = y[i];
```



Vector Assignment (Copy) Kernel

$x \leftarrow y$ for (large) vectors x, y




Parameters (1900 variations)

```
for (size_t i = group_start + get_local_id(0);  
    i < group_end; i += get_local_size(0))  
    x[i] = y[i];
```

Vector Assignment (Copy) Kernel

$x \leftarrow y$ for (large) vectors x, y



Parameters (1900 variations)


```
for (size_t i = group_start + get_local_id(0);  
    i < group_end; i+= get_local_size(0))  
    x[i] = y[i];
```

Benchmark Setting

Operations

Vector copy, vector addition, inner product

Matrix-vector product




Devices


AMD: A10-5800 APU, HD 5850 GPU

INTEL: Dual Socket Xeon E5-2670, Xeon Phi


NVIDIA: GTX 285, Tesla K20m




Histograms




AMD Radeon HD 5850



NVIDIA Tesla K20m




Intel Xeon E5-2670




Bandwidth Inner Product (% of theoretical peak)

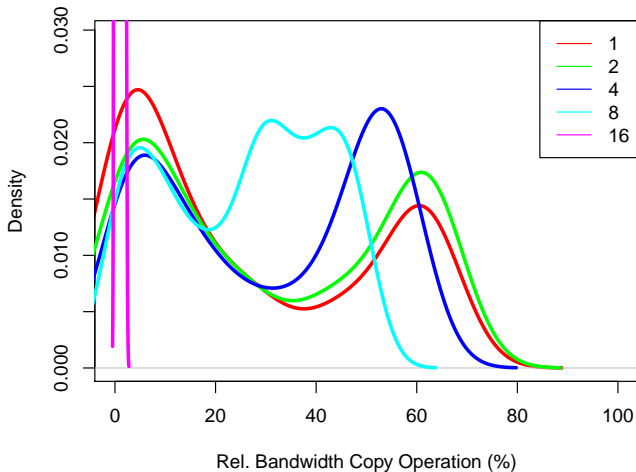
Intel Xeon E5-2670



Bandwidth Inner Product (% of theoretical peak)




NVIDIA Tesla K20m



(comparison of vector types double, double2, double4, double8, double16)

Benchmark

NVIDIA Tesla K20m




Rel. Bandwidth Matrix-Vector Product Operation (%)
(comparison of vector types double, double2, double4, double8, double16)

[Addition|Inner Product|Matrix-Vector] vs. Copy Kernel


Same Device




NVIDIA GeForce GTX 285



NVIDIA GeForce GTX 285




NVIDIA GeForce GTX 285




NVIDIA Tesla K20m


Addition



Inner Product




Mat-Vec Product




AMD Radeon HD 5850


Addition



Inner Product




Mat-Vec Product




INTEL Dual Xeon E5-2670


Addition



Inner Product




Mat-Vec Product




INTEL Xeon Phi


Addition



Inner Product




Mat-Vec Product



Conclusio:

Focus on fastest configurations for copy-kernel sufficient




[Copy|Addition|Inner Product|Matrix-Vector] vs. Copy Kernel


Different Device, Same Vendor




NVIDIA Hardware (x: GTX 285, y: K20m)



AMD Hardware (x: A10-5800K GPU, y: HD 5850)




INTEL Hardware (x: Xeon Phi, E5-2670)




NVIDIA Hardware (x: GTX 285, y: K20m)


Addition



Inner Product




Mat-Vec Product




AMD Hardware (x: A10-5800K GPU, y: HD 5850)


Addition



Inner Product



Mat-Vec Product



Conclusio:

Certain Performance Portability per Vendor




[Copy|Addition|Inner Product|Matrix-Vector] vs. Copy Kernel

Different Device, Different Vendor




Benchmark


x: INTEL CPU, y: AMD CPU



x: INTEL CPU, y: NVIDIA GPU




x: AMD GPU, y: NVIDIA GPU




x: AMD HD 5850, y: NVIDIA K20m


Addition



Inner Product



Mat-Vec Product



Conclusio:

Fast Configurations Across Vendors Exist




Introduction



Conjugate Gradient - Standard Formulation



Conjugate Gradient - Kernel Parameters



Parameter Study for Portable Performance




Conjugate Gradient - Pipelining




GMRES Optimizations




Conclusion



Conjugate Gradients



Conjugate Gradients




Optimization 3: Rearrange the algorithm

- Remove unnecessary reads

- Remove unnecessary synchronizations

- Use custom kernels instead of standard BLAS



Standard CG

Choose x_0

$$p_0 = r_0 = b - Ax_0$$

For $i = 0$ until convergence

1. Compute and store Ap_i
2. Compute $\langle p_i, Ap_i \rangle$
3. $\alpha_i = \langle r_i, r_i \rangle / \langle p_i, Ap_i \rangle$
4. $x_{i+1} = x_i + \alpha_i p_i$
5. $r_{i+1} = r_i - \alpha_i Ap_i$
6. Compute $\langle r_{i+1}, r_{i+1} \rangle$
7. $\beta_i = \langle r_{i+1}, r_{i+1} \rangle / \langle r_i, r_i \rangle$
8. $p_{i+1} = r_{i+1} + \beta_i p_i$

EndFor

Pipelined CG


Choose x_0

$$p_0 = r_0 = b - Ax_0$$


For $i = 1$ until convergence

1. $i = 1$: Compute α_0, β_0, Ap_0
2. $x_i = x_{i-1} + \alpha_{i-1} p_{i-1}$
3. $r_i = r_{i-1} - \alpha_{i-1} Ap_{i-1}$
4. $p_i = r_i + \beta_{i-1} p_{i-1}$
5. Compute and store Ap_i
6. Compute $\langle Ap_i, Ap_i \rangle, \langle p_i, Ap_i \rangle, \langle r_i, r_i \rangle$
7. $\alpha_i = \langle r_i, r_i \rangle / \langle p_i, Ap_i \rangle$
8. $\beta_i = (\alpha_i^2 \langle Ap_i, Ap_i \rangle - \langle r_i, r_i \rangle) / \langle r_i, r_i \rangle$

EndFor




Conjugate Gradients



Conjugate Gradients

Time per CG Iteration - AMD FirePro W9000



Generalized Minimum Residual (GMRES) Method

Krylov space $\text{span}\{r, Ar, A^2r, \dots, A^{N-1}r\}$

Orthogonal basis $\{v_1, v_2, \dots, v_N\}$

Gram-Schmidt Method revisited

Given: orthonormal basis $\{v_1, v_2, \dots, v_N\}$, augment by w

$$w \leftarrow w - \sum_{i=1}^N \langle w, v_i \rangle v_i$$

$$w \leftarrow w / \|w\|$$

Add w to basis

Multiple inner products $\langle w, v_i \rangle$

Performance critical (global reductions)

Reuse of w desirable



Custom routine *mdot*


Process $\alpha_i = \langle w, v_i \rangle$ in batches

Batch sizes 1, 2, 3, 4, 8


Batch size 8: Only 12.5% overhead vs. arbitrary batch sizes

Code sketch (Batch size 4)

```
for (size_t i=thread_id; i<M; i += threads_per_group)
{
    double val_w = w[i];
    alpha_1 += val_w * v1[i];
    alpha_2 += val_w * v2[i];
    alpha_3 += val_w * v3[i];
    alpha_4 += val_w * v4[i];
}
```




Benchmarks



Fixed number of 50 vectors

Benchmarks



Fixed number of 50 vectors

Conjugate Gradient Method

- Careful choice of sparse matrix format

- Tune kernels to target device

- Minimize reads from global memory (kernel fusion, pipelining)

Generalized Minimum Residual Method (GMRES)

- Minimizes reads from global memory (mdot kernel)

- Up to twice the performance of 'naive' implementations

Implications

- Tune primarily with memory transfers in mind

- Prefer regular memory access patterns

- Use *appropriate* vector data types

