

Matrix Multiplications: A Lot of Ways of Handling a Few Nonzeros

Karl Rupp

Short-Term Visiting Scientist
Argonne National Laboratory



extending previous work at TU Wien with
A. Morhammer, F. Rudolf, and J. Weinbub



LANS Informal Seminar
September 27, 2017

Positions

PhD student at TU Wien (2009-2011)

Postdoc at Argonne Natl. Lab. (09/2012-09/2013)

Postdoc at TU Wien (09/2013-10/2015; 01/2017-)

Freelancer, ETH Zürich (03/2016-06/2017)

Research Interests

Numerical solution of PDEs

Semiconductor device simulation

Parallel computing

Software Development

PETSc

ViennaCL

ViennaSHE

...


Matrix Multiplication

Dense Matrix-Matrix-Multiplications

Ubiquitous for: dense linear algebra (eigenvalues, LU factorization, etc.)

FLOP-limited, basis for TOP500

Computer scientist's darling




https://en.wikipedia.org/wiki/Matrix_multiplication#/media/File:Matrix_multiplication_diagram_2.svg

Compressed Sparse Row Format

Sparse Matrix Storage


	2		7	4
		1	9	
3				
		6		5
	8			

Matrix



CSR

Sparse Matrix Products




Sparse Matrix Products



$$\text{Row } i: \quad \mathbf{c}_i = \sum_j a_{ij} \mathbf{b}_j$$

Sparse Matrix Products



$$\text{Row } i: \quad \mathbf{c}_i = \sum_j a_{ij} \mathbf{b}_j$$

Sequential MatMatMult in PETSc

Sorted

Scalable

Scalable, fast

Heap

BHeap

Linked-list condensed

RowMerge – **NEW!**

RowMerge2 – **NEW!**

Sequential MatMatMult in PETSc

Sorted

Scalable

Scalable, fast

Heap

BHeap

Linked-list condensed

RowMerge – NEW!

RowMerge2 – NEW!

Two Stages

Symbolic phase: Determine sparsity pattern

Numeric phase: Compute numerical values, sparsity pattern known

Good News:

Numeric phase is easy!

Sparse Matrix-Matrix Multiplications

Result row



1 2 3 5 6 7

Rows to add



1 2 3 4 5 6 7

Numeric Phase

Sparse Matrix-Matrix Multiplications

Result row




1 2 3 5 6 7

Rows to add



Dense Temporary Row




1 2 3 4 5 6 7

Numeric Phase

Merge directly to dense array


Sparse Matrix-Matrix Multiplications



Numeric Phase

Merge directly to dense array

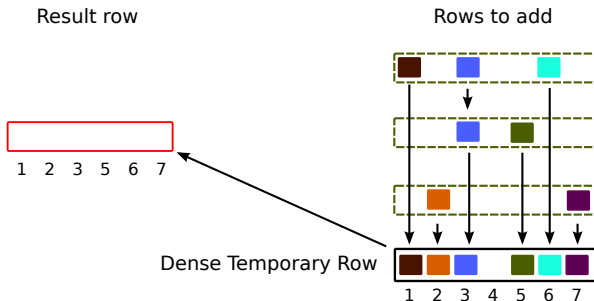
Sparse Matrix-Matrix Multiplications



Numeric Phase

Merge directly to dense array

Sparse Matrix-Matrix Multiplications




Numeric Phase

Merge directly to dense array

Pick up nonzeros to form \mathbf{C}

Sparse Matrix-Matrix Multiplications



Numeric Phase

Merge directly to dense array


Pick up nonzeros to form \mathbf{C}

Bad News:

Symbolic phase is tricky!

Sparse Matrix-Matrix Multiplications

Result row



Rows to add



1 2 3 4 5 6 7


Sorted MatMatMult in PETSc

Dense flag array

Segmented buffer for nonzero indices

Sparse Matrix-Matrix Multiplications

Result row



Rows to add



0	0	0	0	0	0	0
1	2	3	4	5	6	7


Sorted MatMatMult in PETSc

Dense flag array


Segmented buffer for nonzero indices

Sparse Matrix-Matrix Multiplications

Result row



Rows to add



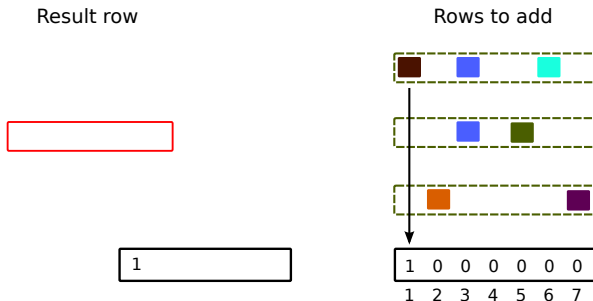
0	0	0	0	0	0	0
1	2	3	4	5	6	7

Sorted MatMatMult in PETSc

Dense flag array

Segmented buffer for nonzero indices

Sparse Matrix-Matrix Multiplications




Sorted MatMatMult in PETSc

Dense flag array

Segmented buffer for nonzero indices


Sparse Matrix-Matrix Multiplications

Result row



1 3

Rows to add




Sorted MatMatMult in PETSc

Dense flag array

Segmented buffer for nonzero indices


Sparse Matrix-Matrix Multiplications

Result row



1 3 6

Rows to add



Sorted MatMatMult in PETSc

Dense flag array

Segmented buffer for nonzero indices


Sparse Matrix-Matrix Multiplications

Result row



1 3 6 5

Rows to add




Sorted MatMatMult in PETSc

Dense flag array

Segmented buffer for nonzero indices


Sparse Matrix-Matrix Multiplications

Result row



1	3	6	5	2
---	---	---	---	---

Rows to add




Sorted MatMatMult in PETSc

Dense flag array

Segmented buffer for nonzero indices


Sparse Matrix-Matrix Multiplications

Result row



1	3	6	5	2	7
---	---	---	---	---	---

Rows to add




Sorted MatMatMult in PETSc

Dense flag array

Segmented buffer for nonzero indices

Sparse Matrix-Matrix Multiplications

Result row




1 2 3 5 6 7

sort
↙

1 3 6 5 2 7

Rows to add



Sorted MatMatMult in PETSc

Dense flag array

Segmented buffer for nonzero indices

Sparse Matrix-Matrix Multiplications

Result row




Rows to add



1 2 3 4 5 6 7

Scalable MatMatMult in PETSc

Sparse Matrix-Matrix Multiplications



Scalable MatMatMult in PETSc

Use {linked list, heap, binary tree} to merge nonzero column indices

No dense array

Sparse Matrix-Matrix Multiplications

Result row



Rows to add



1 2 3 4 5 6 7


Scalable MatMatMult in PETSc

Use {linked list, heap, binary tree} to merge nonzero column indices

No dense array

Sparse Matrix-Matrix Multiplications

Result row



Rows to add



1 2 3 4 5 6 7

Scalable MatMatMult in PETSc

Use {linked list, heap, binary tree} to merge nonzero column indices

No dense array

Sparse Matrix-Matrix Multiplications

Result row



Rows to add



1 2 3 4 5 6 7


Scalable MatMatMult in PETSc

Use {linked list, heap, binary tree} to merge nonzero column indices

No dense array

Sparse Matrix-Matrix Multiplications

Result row



Rows to add




1 2 3 4 5 6 7

Scalable MatMatMult in PETSc

Use {linked list, heap, binary tree} to merge nonzero column indices

No dense array


Sparse Matrix-Matrix Multiplications



Scalable MatMatMult in PETSc

Use {linked list, heap, binary tree} to merge nonzero column indices
No dense array


Sparse Matrix-Matrix Multiplications



Scalable MatMatMult in PETSc

Use {linked list, heap, binary tree} to merge nonzero column indices
No dense array

Sparse Matrix-Matrix Multiplications



Scalable MatMatMult in PETSc


Use {linked list, heap, binary tree} to merge nonzero column indices
No dense array

Row Merge Algorithm


Result row



Front




Rows to add



Result Row Computation

1. Determine minimum index in front


Row Merge Algorithm



Result Row Computation

1. Determine minimum index in front
2. Write minimum index


Row Merge Algorithm



Result Row Computation

1. Determine minimum index in front
2. Write minimum index
3. Advance front where minimum index occurred


Row Merge Algorithm



Result Row Computation

1. Determine minimum index in front
2. Write minimum index
3. Advance front where minimum index occurred


Row Merge Algorithm



Result Row Computation

1. Determine minimum index in front
2. Write minimum index
3. Advance front where minimum index occurred


Row Merge Algorithm



Result Row Computation

1. Determine minimum index in front
2. Write minimum index
3. Advance front where minimum index occurred

Row Merge Algorithm




Result Row Computation

1. Determine minimum index in front
2. Write minimum index
3. Advance front where minimum index occurred

Algorithm Details

Split matrix if rows too large
Recursively merge 8 rows
(Re-)use scratchpad memory




<https://en.wikipedia.org/wiki/Octree\#/media/File:Octree2.svg>

Sparse Matrix Products

Algorithm Details

Split matrix if rows too large
Recursively merge 8 rows
(Re-)use scratchpad memory



<https://en.wikipedia.org/wiki/Octree\#/media/File:Octree2.svg>

Hardware Details

CPU: Use AVX2 to merge 8 rows simultaneously


GPU: Merge up to 32/64 rows simultaneously with one warp/wavefront

MIC: Use AVX-KNC to merge 16 rows simultaneously

Sparse Matrix Products

Algorithm Details

Split matrix if rows too large
Recursively merge 8 rows
(Re-)use scratchpad memory



<https://en.wikipedia.org/wiki/Octree\#/media/File:Octree2.svg>

Hardware Details

CPU: Use AVX2 to merge 8 rows simultaneously

GPU: Merge up to 32/64 rows simultaneously with one warp/wavefront

MIC: Use AVX-KNC to merge 16 rows simultaneously

Parallelization and Load Balancing

CPU, MIC: Use MPI!

CPU, MIC, alternative: OpenMP dynamic scheduling, chunk size 1024


GPU: thread scheduling in hardware

Which approach is fastest?

Sorted vs. Scalable vs. Scalable, fast vs. Heap vs. BTHHeap vs.
Linked-list condensed vs. RowMerge vs. RowMerge2

Benchmark Results

Comparison of Sequential MatMatMult in PETSc



(Tests run on an Intel Core i3-3217U)

Row Merge up to 2x faster than PETSc's default

Hardware for Comparison

AMD FirePro W9100 GPU

Intel Xeon E5-2670v3 (Haswell, dual socket)

NVIDIA Tesla K20

Software for Comparison

Intel MKL 11.2.1

NVIDIA cuSPARSE 7.0


CUSP 0.5.1

Matrices

20 matrices from Florida Sparse Matrix Collection


Operation: $\mathbf{B} = \mathbf{A}\mathbf{A}$

Benchmark Results




Intel Xeon E5-2670v3

Benchmark Results



AMD FirePro W9100, NVIDIA Tesla K20m

Benchmark Results




Comparison on Intel Xeon E5-2670v3 (Haswell) vs. Xeon Phi (KNL)

Sparse Matrix-Matrix Multiplications


Parallel MatMatMult in PETSc

Scalable


Non-Scalable



C



A




B

Sparse Matrix-Matrix Multiplications


Parallel MatMatMult in PETSc

Scalable

Non-Scalable



C



A


B

Sparse Matrix-Matrix Multiplications

Parallel MatMatMult in PETSc

Scalable


Non-Scalable



C



A




B

Sparse Matrix-Matrix Multiplications

Parallel MatMatMult in PETSc

Scalable


Non-Scalable



C



A



B

Lots of bookkeeping...

Sequential Sparse Matrix-Matrix Products

Row Merge faster than MKL on Xeon CPUs on average

Row Merge faster than cuSPARSE and CUSP on Tesla K20m (and others)

Sequential Sparse Matrix-Matrix Products

Row Merge faster than MKL on Xeon CPUs on average

Row Merge faster than cuSPARSE and CUSP on Tesla K20m (and others)

Parallel Sparse Matrix-Matrix Products

Potential for an approach similar to Row Merge

Fine-tune use of scalable vs. non-scalable routines

Sequential Sparse Matrix-Matrix Products

- Row Merge faster than MKL on Xeon CPUs on average

- Row Merge faster than cuSPARSE and CUSP on Tesla K20m (and others)

Parallel Sparse Matrix-Matrix Products

- Potential for an approach similar to Row Merge

- Fine-tune use of scalable vs. non-scalable routines

Implications and Outlook

- CPUs beat accelerators for sparse matrix-matrix products (caches!)

- Full integration into PETSc's GAMG (algebraic multigrid)

- Find CPU/GPU balance for hybrid clusters