

Frequently and Not-So-Frequently Asked PETSc-Questions

Karl Rupp

`me@karlrupp.net`

Freelance Computational Scientist
and
Institute for Microelectronics, TU Wien



PETSc User Meeting 2017

Boulder, Colorado

June 14-16, 2017

What about threads in PETSc?

What about all those branches in the repository?

What kind of questions do users have?

What is Firetran?

What about threads in PETSc?

Competing Threading Approaches

pthread

OpenMP

C++11 threads

Compiler magic

Issues with Threads

Problematic across compilers

Data locality

Thread ownership

Software interface

Assumption for Subsequent Discussion

Primary Goal: Get the science done!


10 percent performance difference is **not significant**

When to Use Threads After All?

Distributed Memory: Need MPI anyway

Shared Memory: MPI for multi-socket systems for NUMA reasons

1-Socket Machines: If a 2-5x gain is critical, use a cluster!



Attempt 1

Library spawns threads

```
void library_func(double *x, int N) {  
    #pragma omp parallel for  
    for (int i=0; i<N; ++i) x[i] = something_complicated();  
}
```

Problems

Call from multi-threaded environment?

```
void user_func(double **y, int N) {  
    #pragma omp parallel for  
    for (int j=0; j<M; ++j) library_func(y[j], N);  
}
```

Incompatible OpenMP runtimes (e.g. GCC vs. ICC)

Attempt 2

Use pthreads/TBB/etc. instead of OpenMP to spawn threads
Fixes incompatible OpenMP implementations (probably)

Problems

Still a problem with multi-threaded user environments

```
void user_func(double **y, int N) {  
    #pragma omp parallel for  
    for (int j=0; j<M; ++j) library_func(y[j], N);  
}
```

Attempt 3

Hand back thread management to user

```
void library_func(ThreadInfo ti, double *x, int N) {  
    int start = compute_start_index(ti, N);  
    int stop  = compute_stop_index(ti, N);  
    for (int i=start; i<stop; ++i)  
        x[i] = something_complicated();  
}
```

Implications

Users can use their favorite threading model

API requires one extra parameter

Extra boilerplate code required in user code

Reflection

Extra thread communication parameter

```
void library_func(ThreadInfo ti, double *x, int N) {...}
```

Rename thread management parameter

```
void library_func(Thread_Comm c, double *x, int N) {...}
```

Compare:

```
void library_func(MPI_Comm comm, double *x, int N) {...}
```

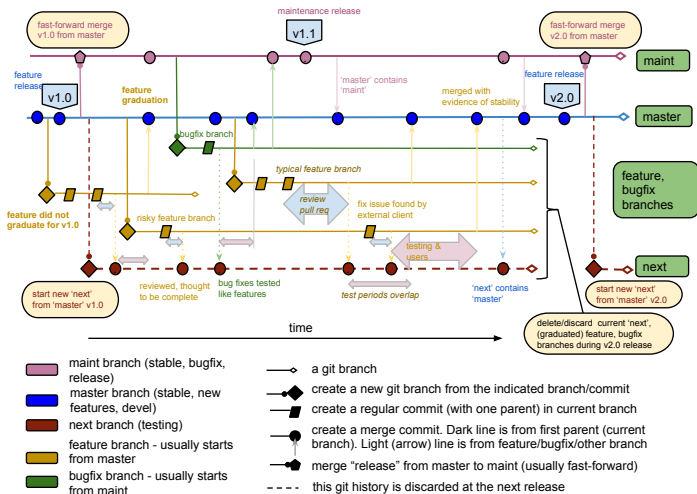
Conclusion

Prefer flat MPI over MPI+OpenMP for a composable software stack


MPI automatically brings better data locality

What about all those branches in the repository?

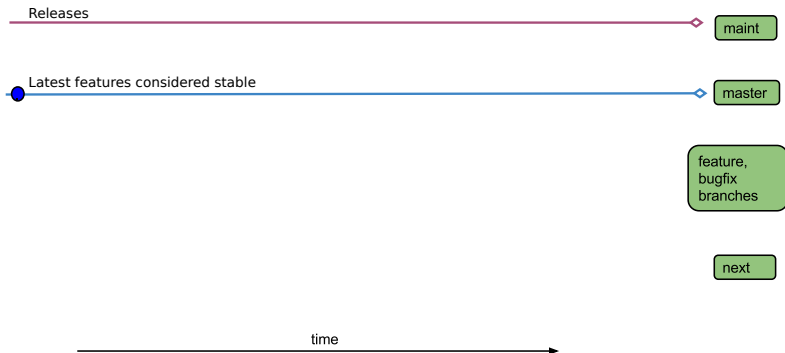
PETSc's Workflow



PETSc and Git




PETSc and Git




PETSc and Git



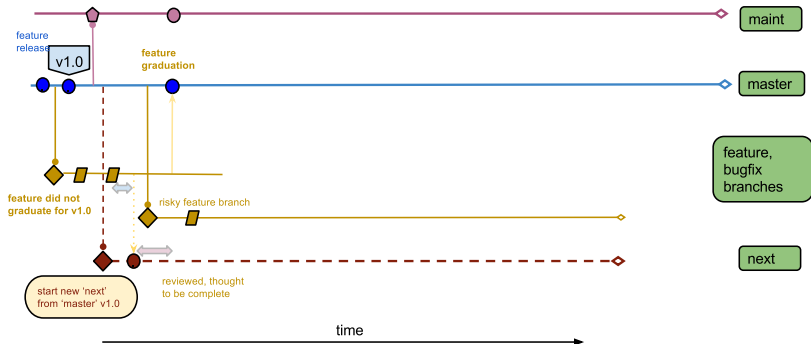
PETSc and Git



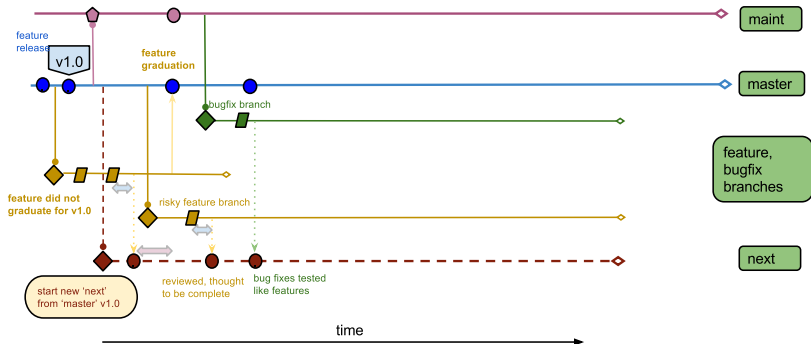
PETSc and Git



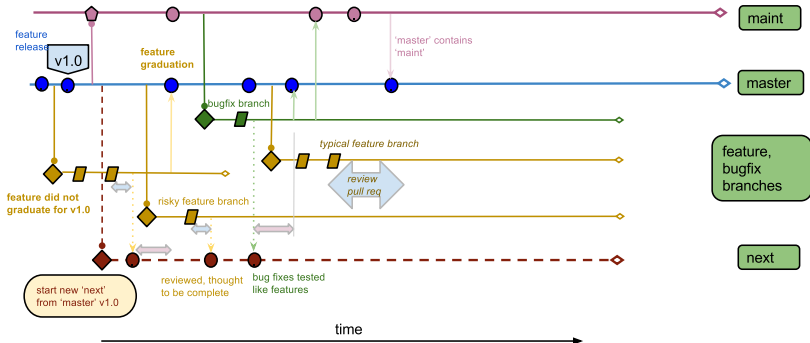
PETSc and Git



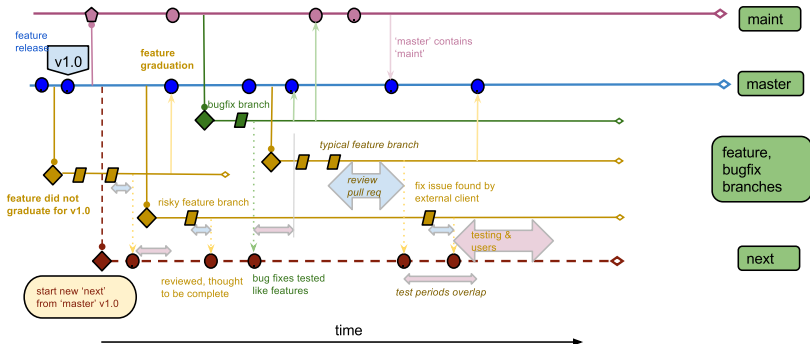
PETSc and Git



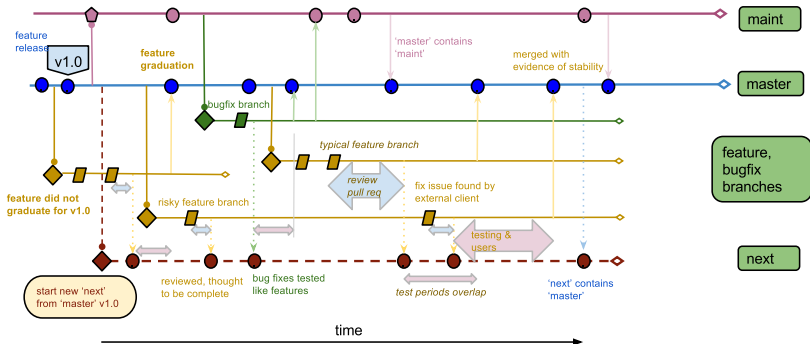
PETSc and Git



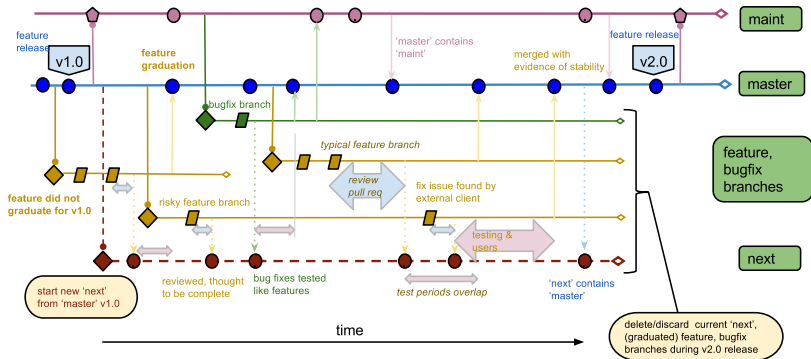
PETSc and Git



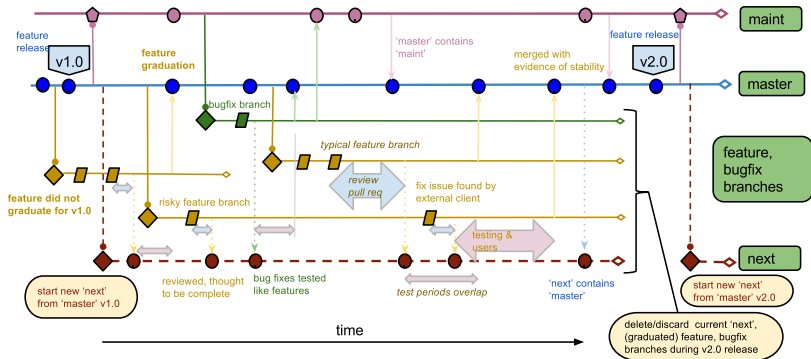
PETSc and Git



PETSc and Git



PETSc and Git




What kind of questions do users have?


Study Setup

Classification of Email threads for petsc-user

Years 2006 - 2014

Volume: 904 emails in 2006, 3947 emails in 2014





Conclusion

Involvement of core developers essential

No significant shift in type of support emails

“Better” documentation would not change email volume significantly

What is Firetran?

Firetran

My New Web Browser

after J. Brown, M. Knepley, B. Smith: "Run-time extensibility and librarization of simulation software" (2015)

Performance

- 10 percent faster than Firefox ...

- ... but no JavaScript

- Recompile for JavaScript, performance not shown

- Character encoding compiled in for efficiency

Plugins

- Directly coded in the core, guarded by `#ifdef`

- Some ship their own version of Firetran

Code Quality

- Some send bug reports

- Ignore if not reproducible with my version

Convenience

Proxy configuration compiled in, don't worry about runtime dialogs
To change: Just edit makefile and recompile

Security

Firetran speaks either http or https
Firetran-http and Firetran-https

Open Source

Development done in private
Release when paper comes out
“We fixed that in our private repository last year”
Fill out attached form and fax signed copy to view our repository

Parental Filter

up to 16 pages compiled in

Building Firetran

Build with last year's version of ACME Fortran 77

Build system consists of csh, perl, m4, BSD make

Navigation

No URL bar to save precious screen space

URL entry in configuration file, restart Firetran

Student wrote Tcl script to automate editing and restarting

Script hard to understand, but the way forward

Only One Problem

Firetran struggles with market share

Reasons

Laughably unacceptable for non-scientific software

Some Linear Algebra and Solver Libraries for Many-Core

MAGMA, cMAGMA, MAGMA MIC: CUDA, OpenCL, OpenMP/AVX

MKL: Web tool for linker flags

SuperLU, SuperLU_dist: Shared vs. distributed memory

Paralution, VexCL: Backend selection at compiletime

ViennaCL: Enable backends at compiletime

(long list to continue...)

Possible Approaches to Fix

Raise awareness for usability issues

Incentives for better software engineering

Don't penalize scientists for writing reusable software

PETSc can Help You

- Solve algebraic and DAE problems in your application area
- Rapidly develop efficient parallel code, can start from examples
- Develop new solution methods and data structures
- Debug and analyze performance
- Advice on software design, solution algorithms, and performance

`petsc-{users,dev,maint}@mcs.anl.gov`

You can Help PETSc

- Report bugs and inconsistencies, or if you think there is a better way
- Tell us if the documentation is inconsistent or unclear
- Consider developing new algebraic methods as plugins, contribute if your idea works

Thank you!