

PETSc Tutorial

PETSc and Many-Core/GPU Architectures

Karl Rupp

me@karlrupp.net

Freelance Computational Scientist
and
Institute for Microelectronics, TU Wien

Boulder, Colorado

June 14-16, 2017



FLOPs and Bandwidth

Performance Modeling

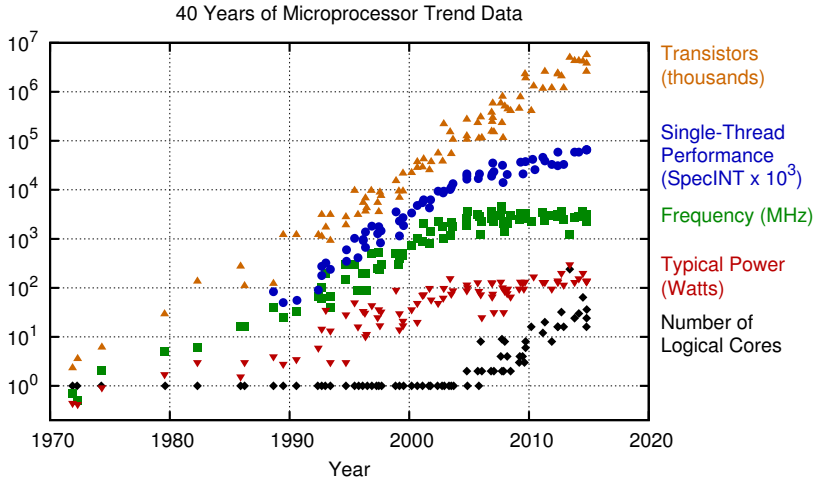
PETSc Profiling

PETSc and Threads

PETSc and GPUs

FLOPs and Bandwidth

Introduction

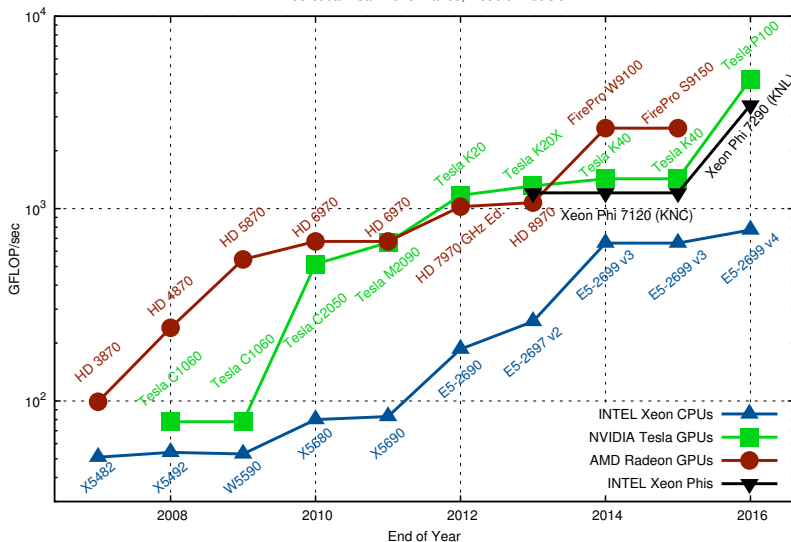


Original data up to the year 2010 collected and plotted by M. Horowitz, F. Labonte, O. Shacham, K. Olukotun, L. Hammond, and C. Batten
New plot and data collected for 2010-2015 by K. Rupp

Introduction

Theoretical Peak Performance

Theoretical Peak Performance, Double Precision

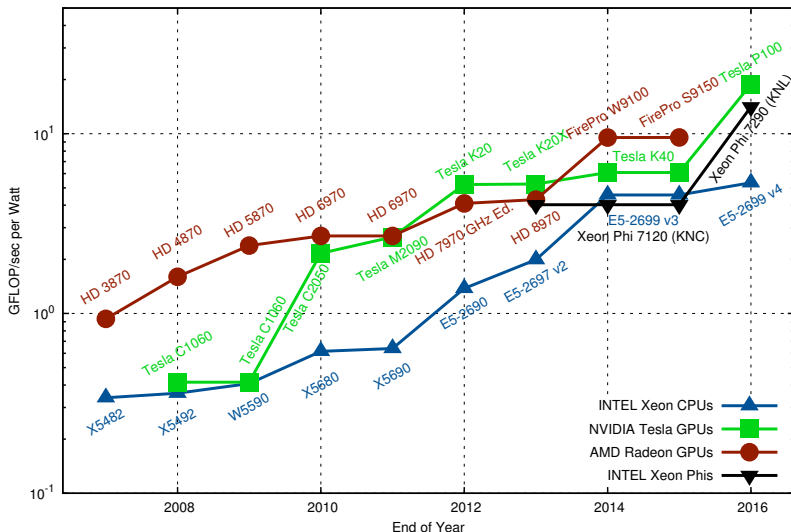


<https://www.karlsruhp.net/2013/06/cpu-gpu-and-mic-hardware-characteristics-over-time/>

Introduction

Theoretical Peak Performance per Watt

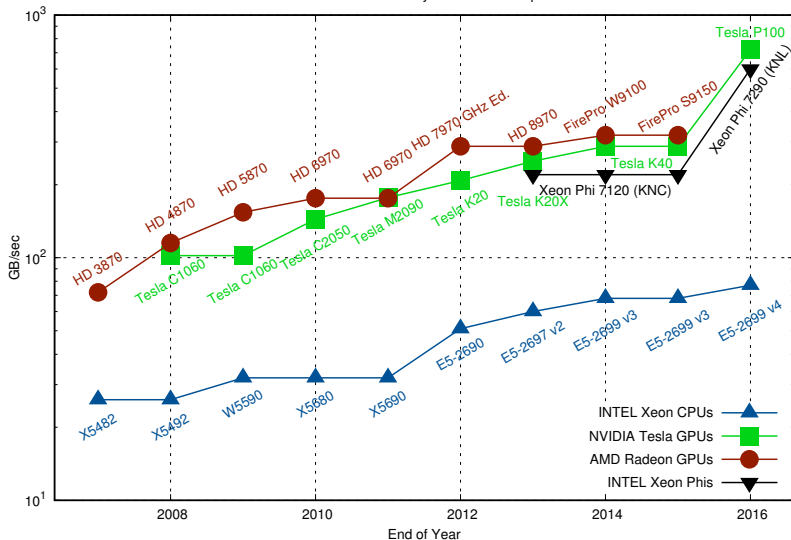
Theoretical Peak Floating Point Operations per Watt, Double Precision



Introduction

Memory Bandwidth

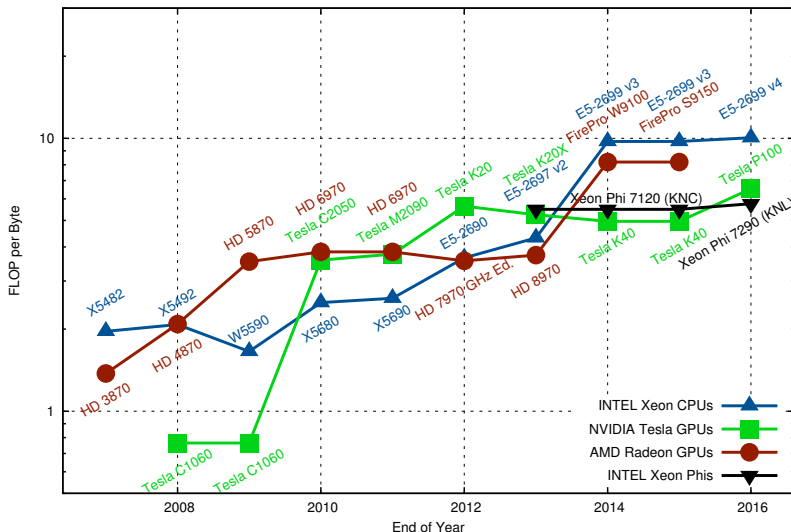
Theoretical Peak Memory Bandwidth Comparison



Introduction

Theoretical Peak Performance (FLOPs) per Byte of Memory Bandwidth

Theoretical Peak Floating Point Operations per Byte, Double Precision



<https://www.karlrupp.net/2013/06/cpu-gpu-and-mic-hardware-characteristics-over-time/>

Typical PETSc Operations

- Vector operations (add, dot, etc.)

- Sparse matrix-vector products (Krylov solvers, smoothers, residuals, etc.)

Maximizing Memory Bandwidth

- Read contiguous blocks of memory (contiguous access)

- Avoid unordered reads whenever possible

Check Memory Bandwidth Yourself

make streams

Performance question to `petsc-maint`, about 2h ago:

```
np  speedup
```

```
1   1.0
```

```
2   1.85
```

```
3   2.25
```

```
4   2.37
```

```
(...)
```

```
39  2.47
```

```
40  2.45
```

Estimation of possible speedup of MPI programs based on
Streams benchmark.

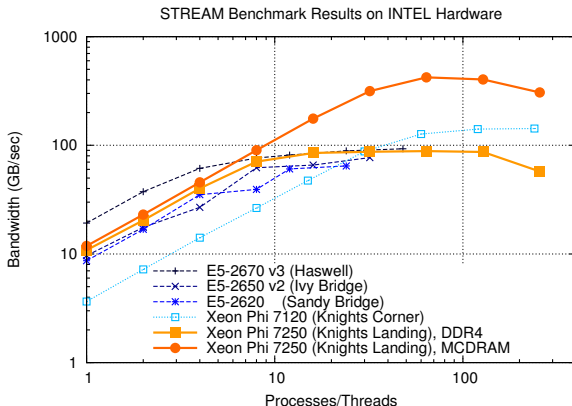
It appears you have 1 node(s)

FLOPs and Bandwidth

How does Memory Bandwidth Scale with Cores?

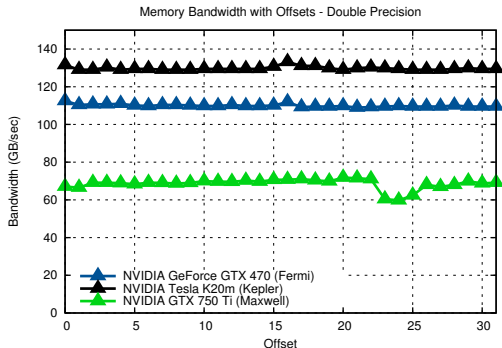
Usually saturates quickly

8-16 processes/threads usually suffice



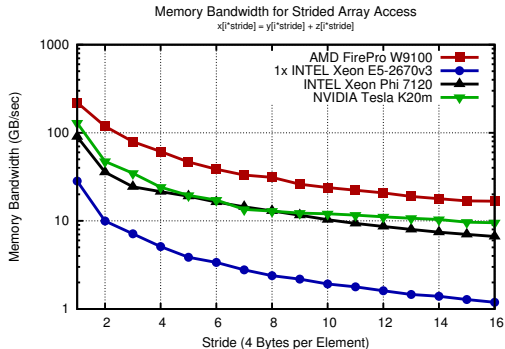
Offset Memory Access

```
void work(double *x, double *y, double *z, int N, int k)
{
    for (size_t i=0; i<N; ++i)
        z[i+k] = x[i+k] + y[i+k];
}
```



Strided Memory Access

```
void work(double *x, double *y, double *z, int N, int k)
{
    for (size_t i=0; i<N; ++i)
        z[i*k] = x[i*k] + y[i*k];
}
```

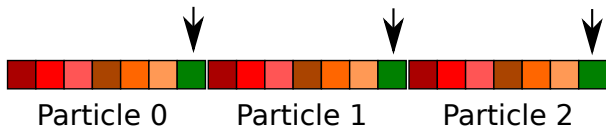


Strided Memory Access

Array of structs problematic

```
typedef struct particle
{
    double pos_x; double pos_y; double pos_z;
    double vel_x; double vel_y; double vel_z;
    double mass;
} Particle;

void increase_mass(Particle *particles, int N)
{
    for (int i=0; i<N; ++i)
        particles[i].mass *= 2.0;
}
```



Strided Memory Access

Workaround: Structure of Arrays

```
typedef struct particles
{
    double *pos_x; double *pos_y; double *pos_z;
    double *vel_x; double *vel_y; double *vel_z;
    double *mass;
} Particle;

void increase_mass(Particle *particles, int N)
{
    for (int i=0; i<N; ++i)
        particles.mass[i] *= 2.0;
}
```



**Mindlessly applying object-oriented programming
all the way down to fine granularity
is a recipe for a performance disaster.**

Performance Modeling

*At any given time during a run,
at least one hardware component
is operating at 100 percent capacity.*

Latency

Bottleneck in strong scaling limit

Ultimate limit for time stepping

Latency - Sources

Network latency (Ethernet $\sim 20\mu\text{s}$, Infiniband $\sim 5\mu\text{s}$)

PCI-Express latency (Kernel launches, $\sim 10\mu\text{s}$)

Thread synchronization (barriers, locks, $\sim 1 - 100\mu\text{s}$)

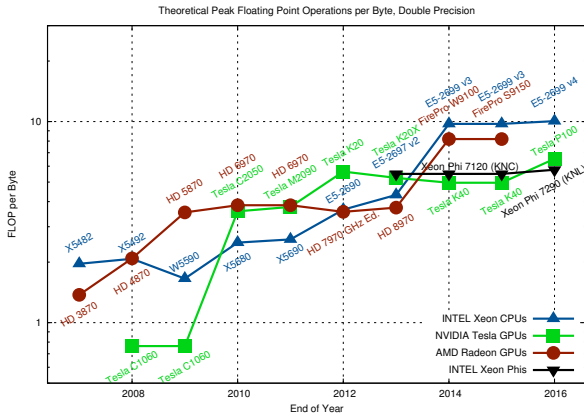
Memory latency ($\sim 100\text{ns}$)

Arithmetic Intensity

Number of FLOPs per Byte

FLOP-limited: Arithmetic intensity larger than ~ 10

Memory-limited: Arithmetic intensity smaller than ~ 1



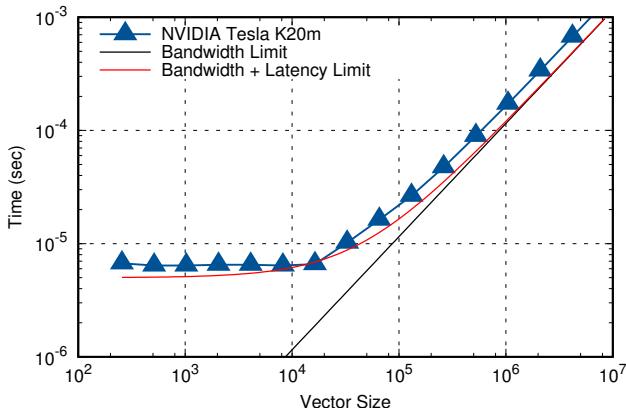
Vector Addition

$x = y + z$ with N elements each

1 FLOP per 24 byte in double precision

Limited by memory bandwidth $\Rightarrow T_2(N) \stackrel{?}{\approx} 3 \times 8 \times N / \text{Bandwidth} + \text{Latency}$

Execution Time for $x = y + z$ in Double Precision



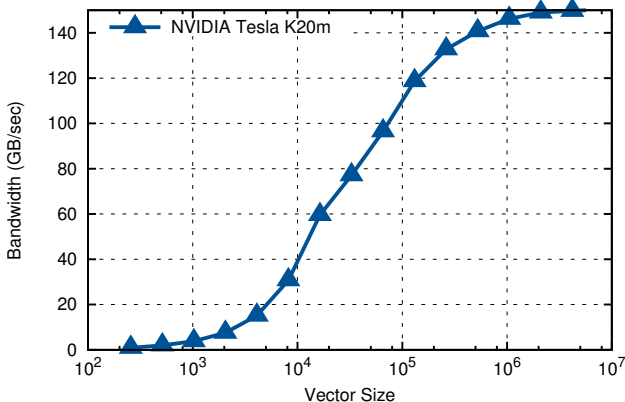
Vector Addition

$x = y + z$ with N elements each

1 FLOP per 24 byte in double precision

Limited by memory bandwidth $\Rightarrow T_2(N) \stackrel{?}{\approx} 3 \times 8 \times N / \text{Bandwidth} + \text{Latency}$

Memory Bandwidth for $x = y + z$ in Double Precision



PETSc Profiling

First: Get the Math Right!

Choose an algorithm that gives robust iteration counts

Choose an algorithm that really converges

Profiling

Use `-log_view` for a performance profile

- Event timing

- Event flops

- Memory usage

- MPI messages

Call `PetscLogStagePush()` and `PetscLogStagePop()`

- User can add new stages

Call `PetscLogEventBegin()` and `PetscLogEventEnd()`

- User can add new events

Call `PetscLogFlops()` to include your flops

Reading -log_view

	Max	Max/Min	Avg	Total
Time (sec):	1.548e+02	1.00122	1.547e+02	
Objects:	1.028e+03	1.00000	1.028e+03	
Flops:	1.519e+10	1.01953	1.505e+10	1.204e+11
Flops/sec:	9.814e+07	1.01829	9.727e+07	7.782e+08
MPI Messages:	8.854e+03	1.00556	8.819e+03	7.055e+04
MPI Message Lengths:	1.936e+08	1.00950	2.185e+04	1.541e+09
MPI Reductions:	2.799e+03	1.00000		

Also a summary per stage

Memory usage per stage (based on when it was allocated)

Time, messages, reductions, balance, flops per event per stage

Always send `-log_view` when asking
performance questions on mailing list

PETSc Profiling

Event	Count		Time (sec)		Flops					--- Global ---					--- Stage ---				
	Max	Ratio	Max	Ratio	Max	Ratio	Mess	Avg len	Reduct	%T	%F	%M	%L	%R	%T	%F	%M	%L	%R
--- Event Stage 1: Full solve																			
VecDot	43	1.0	4.8879e-02	8.3	1.77e+06	1.0	0.0e+00	0.0e+00	4.3e+01	0	0	0	0	0	0	0	0	0	1
VecMDot	1747	1.0	1.3021e+00	4.6	8.16e+07	1.0	0.0e+00	0.0e+00	1.7e+03	0	1	0	0	14	1	1	0	0	27
VecNorm	3972	1.0	1.5460e+00	2.5	8.48e+07	1.0	0.0e+00	0.0e+00	4.0e+03	0	1	0	0	31	1	1	0	0	61
VecScale	3261	1.0	1.6703e-01	1.0	3.38e+07	1.0	0.0e+00	0.0e+00	0.0e+00	0	0	0	0	0	0	0	0	0	0
VecScatterBegin	4503	1.0	4.0440e-01	1.0	0.00e+00	0.0	6.1e+07	2.0e+03	0.0e+00	0	0	50	26	0	0	0	96	53	0
VecScatterEnd	4503	1.0	2.8207e+00	6.4	0.00e+00	0.0	0.0e+00	0.0e+00	0.0e+00	0	0	0	0	0	0	0	0	0	0
MatMult	3001	1.0	3.2634e+01	1.1	3.68e+09	1.1	4.9e+07	2.3e+03	0.0e+00	11	22	40	24	0	22	44	78	49	0
MatMultAdd	604	1.0	6.0195e-01	1.0	5.66e+07	1.0	3.7e+06	1.3e+02	0.0e+00	0	0	3	0	0	0	1	6	0	0
MatMultTranspose	676	1.0	1.3220e+00	1.6	6.50e+07	1.0	4.2e+06	1.4e+02	0.0e+00	0	0	3	0	0	1	1	7	0	0
MatSolve	3020	1.0	2.5957e+01	1.0	3.25e+09	1.0	0.0e+00	0.0e+00	0.0e+00	9	21	0	0	0	18	41	0	0	0
MatCholFctrSym	3	1.0	2.8324e-04	1.0	0.00e+00	0.0	0.0e+00	0.0e+00	0.0e+00	0	0	0	0	0	0	0	0	0	0
MatCholFctrNum	69	1.0	5.7241e+00	1.0	6.75e+08	1.0	0.0e+00	0.0e+00	0.0e+00	2	4	0	0	0	4	9	0	0	0
MatAssemblyBegin	119	1.0	2.8250e+00	1.5	0.00e+00	0.0	2.1e+06	5.4e+04	3.1e+02	1	0	2	24	2	2	0	3	47	5
MatAssemblyEnd	119	1.0	1.9689e+00	1.4	0.00e+00	0.0	2.8e+05	1.3e+03	6.8e+01	1	0	0	0	1	1	0	0	0	1
SNESolve	4	1.0	1.4302e+02	1.0	8.11e+09	1.0	6.3e+07	3.8e+03	6.3e+03	51	50	52	50	50	99100	99100	97		
SNESLineSearch	43	1.0	1.5116e+01	1.0	1.05e+08	1.1	2.4e+06	3.6e+03	1.8e+02	5	1	2	2	1	10	1	4	4	3
SNESFunctionEval	55	1.0	1.4930e+01	1.0	0.00e+00	0.0	1.8e+06	3.3e+03	8.0e+00	5	0	1	1	0	10	0	3	3	0
SNESJacobianEval	43	1.0	3.7077e+01	1.0	7.77e+06	1.0	4.3e+06	2.6e+04	3.0e+02	13	0	4	24	2	26	0	7	48	5
KSPGMRESOrthog	1747	1.0	1.5737e+00	2.9	1.63e+08	1.0	0.0e+00	0.0e+00	1.7e+03	1	1	0	0	14	1	2	0	0	27
KSPSetup	224	1.0	2.1040e-02	1.0	0.00e+00	0.0	0.0e+00	0.0e+00	3.0e+01	0	0	0	0	0	0	0	0	0	0
KSPSolve	43	1.0	8.9988e+01	1.0	7.99e+09	1.0	5.6e+07	2.0e+03	5.8e+03	32	49	46	24	46	62	99	88	48	88
PCSetup	112	1.0	1.7354e+01	1.0	6.75e+08	1.0	0.0e+00	0.0e+00	8.7e+01	6	4	0	0	1	12	9	0	0	1
PCSetupOnBlocks	1208	1.0	5.8182e+00	1.0	6.75e+08	1.0	0.0e+00	0.0e+00	8.7e+01	2	4	0	0	1	4	9	0	0	1
PCApply	276	1.0	7.1497e+01	1.0	7.14e+09	1.0	5.2e+07	1.8e+03	5.1e+03	25	44	42	20	41	49	88	81	39	79

Communication Costs

Reductions: usually part of Krylov method, latency limited

VecDot

VecMDot

VecNorm

MatAssemblyBegin

Change algorithm (e.g. IBCGS)

Point-to-point (nearest neighbor), latency or bandwidth

VecScatter

MatMult

PCApply

MatAssembly

SNESFunctionEval

SNESJacobianEval

Compute subdomain boundary fluxes redundantly

Ghost exchange for all fields at once

Better partition

Adding a Logging Event (C)

```
PetscLogEvent  USER_EVENT;
PetscClassId   classid;
PetscLogDouble user_event_flops;

PetscClassIdRegister("class name",&classid);
PetscLogEventRegister("user event",classid,&USER_EVENT);

PetscLogEventBegin(USER_EVENT,0,0,0,0);
    /* code segment to monitor */
PetscLogFlops(user_event_flops);
PetscLogEventEnd(USER_EVENT,0,0,0,0);
```

Adding a Logging Event (Python)

```
with PETSc.logEvent('Reconstruction') as recEvent:
    # All operations are timed in recEvent
    reconstruct(sol)
    # Flops are logged to recEvent
    PETSc.Log.logFlops(user_event_flops)
```

Adding a Logging Stage (C)

```
PetscLogStage stage;  
  
PetscLogStageRegister("name", &stage);  
PetscLogStagePush(stage);  
  
/* Code to Monitor */  
  
PetscLogStagePop();
```

PETSc Profiling

Event	Count		Time (sec)		Flops		Mess	Avg len	Reduct	--- Global ---					--- Stage ---			
	Max	Ratio	Max	Ratio	Max	Ratio				%T	%F	%M	%L	%R	%T	%F	%M	%L
--- Event Stage 0: Main Stage																		
MatMult	178	1.0	7.8040e+01	1.0	2.59e+11	1.0	4.4e+02	2.0e+05	0.0e+00	33	41	6	11	0	51	89	20	24
MatPtAP	10	1.0	2.4870e+01	1.0	5.45e+09	1.0	2.1e+02	3.1e+05	1.8e+02	10	1	3	8	1	16	2	9	18
MatPtAPSymbolic	10	1.0	1.8828e+01	1.0	0.00e+00	0.0	1.2e+02	2.7e+05	8.2e+01	8	0	2	4	0	12	0	5	9
MatPtAPNumeric	10	1.0	6.0428e+00	1.0	5.45e+09	1.0	9.4e+01	3.7e+05	1.0e+02	3	1	1	4	0	4	2	4	9
SNESolve	2	1.0	1.9059e+02	1.0	6.22e+11	1.0	6.6e+03	9.3e+04	3.4e+03	79	99	92	75	16	123213292168			
KSPSolve	2	1.0	1.8230e+02	1.0	6.07e+11	1.0	6.5e+03	9.1e+04	3.2e+03	76	97	89	72	15	118208285161			
PCSetUp	8	1.0	1.6138e+01	1.0	4.81e+09	1.1	1.2e+03	8.1e+04	2.5e+03	7	1	17	12	11	10	2	55	28
PCApply	46	1.0	1.2586e+02	1.0	4.43e+11	1.0	6.3e+03	8.5e+04	2.7e+03	52	70	87	65	12	81152277146			
KSPSolve_FS_0	46	1.0	1.0038e+02	1.0	3.42e+11	1.0	6.2e+03	8.2e+04	2.6e+03	42	54	86	62	12	65117273138			
(...)																		
--- Event Stage 1: MG Apply																		
MatMultMFA11	296	1.0	4.3461e+01	1.0	2.82e+11	1.0	1.2e+03	3.0e+05	0.0e+00	18	45	16	43	0	51	84	24	78
KSPSolve	230	1.0	7.2581e+01	1.0	2.87e+11	1.0	4.5e+03	8.5e+04	2.6e+02	30	46	62	47	1	85	85	91	841
PCApply	642	1.0	1.0269e+01	1.0	1.40e+10	1.1	3.0e+03	8.7e+03	1.8e+02	4	2	42	3	1	12	4	61	6
MGSmooth Level 0	46	1.0	7.8169e+00	1.0	1.06e+10	1.1	3.0e+03	8.3e+03	1.7e+02	3	2	41	3	1	9	3	61	5
MGSmooth Level 1	92	1.0	2.4177e+01	1.0	3.17e+10	1.0	5.0e+02	1.2e+05	4.6e+01	10	5	7	7	0	28	9	10	13
MGResid Level 1	46	1.0	4.3231e+00	1.0	5.77e+09	1.0	9.2e+01	1.2e+05	0.0e+00	2	1	1	1	0	5	2	2	2
MGInterp Level 1	92	1.0	3.5063e-01	1.1	1.09e+08	1.0	9.2e+01	1.5e+04	0.0e+00	0	0	1	0	0	0	0	2	0
MGSmooth Level 2	92	1.0	4.0886e+01	1.0	2.44e+11	1.0	1.0e+03	3.0e+05	4.6e+01	17	39	14	37	0	48	73	20	66
MGResid Level 2	46	1.0	6.8277e+00	1.0	4.39e+10	1.0	1.8e+02	3.0e+05	0.0e+00	3	7	3	7	0	8	13	4	12
MGInterp Level 2	92	1.0	1.0898e+00	1.4	8.47e+08	1.0	9.2e+01	3.8e+04	0.0e+00	0	0	1	0	0	1	0	2	1
(...)																		

PETSc and Threads

Competing Threading Approaches

pthread

OpenMP

C++11 threads

Compiler magic

...

Issues with Threads

Problematic across compilers

Data locality

Thread ownership

Software interface

Assumption for Subsequent Discussion

Primary Goal: Get the science done!

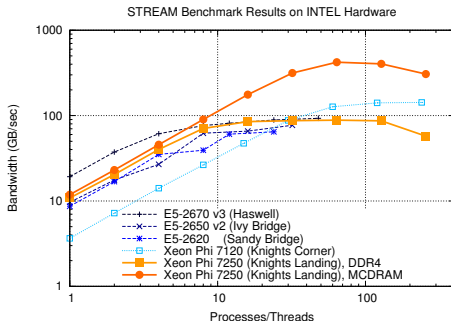
10 percent performance difference is **not significant**

When to Use Threads After All?

Distributed Memory: Need MPI anyway

Shared Memory: MPI for multi-socket systems for NUMA reasons

1-Socket Machines: If a 2-5x gain is critical, use a cluster!



```
numactl --membind 0 ./stream # DDR4  
numactl --membind 1 ./stream # MCDRAM
```

Attempt 1

Library spawns threads

```
void library_func(double *x, int N) {  
    #pragma omp parallel for  
    for (int i=0; i<N; ++i) x[i] = something_complicated();  
}
```

Problems

Call from multi-threaded environment?

```
void user_func(double **y, int N) {  
    #pragma omp parallel for  
    for (int j=0; j<M; ++j) library_func(y[j], N);  
}
```

Incompatible OpenMP runtimes (e.g. GCC vs. ICC)

Attempt 2

Use pthreads/TBB/etc. instead of OpenMP to spawn threads
Fixes incompatible OpenMP implementations (probably)

Problems

Still a problem with multi-threaded user environments

```
void user_func(double **y, int N) {  
    #pragma omp parallel for  
    for (int j=0; j<M; ++j) library_func(y[j], N);  
}
```

Attempt 3

Hand back thread management to user

```
void library_func(ThreadInfo ti, double *x, int N) {  
    int start = compute_start_index(ti, N);  
    int stop  = compute_stop_index(ti, N);  
    for (int i=start; i<stop; ++i)  
        x[i] = something_complicated();  
}
```

Implications

Users can use their favorite threading model

API requires one extra parameter

Extra boilerplate code required in user code

Reflection

Extra thread communication parameter

```
void library_func(ThreadInfo ti, double *x, int N) {...}
```

Rename thread management parameter

```
void library_func(Thread_Comm c, double *x, int N) {...}
```

Compare:

```
void library_func(MPI_Comm comm, double *x, int N) {...}
```

Conclusion

Prefer flat MPI over MPI+OpenMP for a composable software stack

MPI automatically brings better data locality

PETSc and GPUs

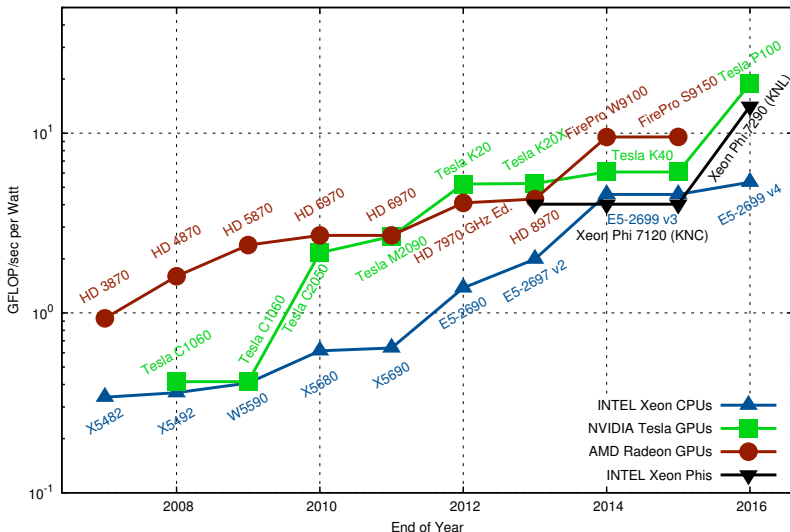
*Don't believe anything
unless you can run it*

Matt Knepley

Why bother?

GFLOPs/Watt

Theoretical Peak Floating Point Operations per Watt, Double Precision



Why bother?

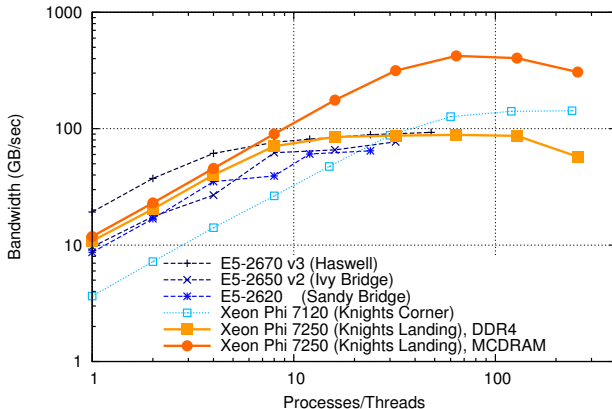
Procurements

Theta (ANL, 2016): 2nd generation INTEL Xeon Phi

Summit (ORNL, 2017), Sierra (LLNL, 2017): NVIDIA Volta GPU

Aurora (ANL, 2018): 3rd generation INTEL Xeon Phi

STREAM Benchmark Results on INTEL Hardware



PETSc on GPUs and MIC:

Current Status

Native on Xeon Phi

Cross-compile for Xeon Phi



CUDA

CUDA-support through CUSP

```
-vec_type cusp -mat_type aijsusp
```

```
-vec_type cuda -mat_type aijsusp
```

Only for NVIDIA GPUs



CUDA/OpenCL/OpenMP

CUDA/OpenCL/OpenMP-support through ViennaCL

```
-vec_type viennacl -mat_type aijsviennacl
```

OpenCL on CPUs and MIC fairly poor



CUDA (CUSP)

CUDA-enabled configuration (minimum)

```
./configure [...] --with-cuda=1  
--with-cusp=1 --with-cusp-dir=/path/to/cusp
```

Customization:

```
--with-cudac=/path/to/cuda/bin/nvcc  
--with-cuda-arch=sm_60
```

OpenCL (ViennaCL)

OpenCL-enabled configuration

```
./configure [...] --download-viennacl  
--with-opencl-include=/path/to/OpenCL/include  
--with-opencl-lib=/path/to/libOpenCL.so
```

How Does It Work?

Host and Device Data

```
struct _p_Vec {  
    ...  
    void          *data;           // host buffer  
    PetscCUSEPFlag valid_GPU_array; // flag  
    void          *spptr;          // device buffer  
};
```

Possible Flag States

```
typedef enum {PETSC_CUSP_UNALLOCATED,  
             PETSC_CUSP_GPU,  
             PETSC_CUSP_CPU,  
             PETSC_CUSP_BOTH} PetscCUSEPFlag;
```

Fallback-Operations on Host

Data becomes valid on host (PETSC_CUSP_CPU)

```
PetscErrorCode VecSetRandom_SeqCUSP_Private(...) {  
    VecGetArray(...);  
    // some operation on host memory  
    VecRestoreArray(...);  
}
```

Accelerated Operations on Device

Data becomes valid on device (PETSC_CUSP_GPU)

```
PetscErrorCode VecAYPX_SeqCUSP(...) {  
    VecCUSPGetArrayReadWrite(...);  
    // some operation on raw handles on device  
    VecCUSPRestoreArrayReadWrite(...);  
}
```

KSP ex12 on Host

```
$> ./ex12  
-pc_type ilu -m 200 -n 200 -log_summary
```

```
KSPGMRESOrthog      228 1.0 6.2901e-01  
KSPSolve            1 1.0 2.7332e+00
```

KSP ex12 on Device

```
$> ./ex12 -vec_type cusp -mat_type aijcusp  
-pc_type ilu -m 200 -n 200 -log_summary
```

```
[0]PETSC ERROR: MatSolverPackage petsc does not support  
matrix type seqaijcusp
```

KSP ex12 on Host

```
$> ./ex12  
      -pc_type none -m 200 -n 200 -log_summary
```

KSPGMRESOrthog	1630	1.0	4.5866e+00
KSPSolve	1	1.0	1.6361e+01

KSP ex12 on Device

```
$> ./ex12 -vec_type cusp -mat_type aijcusp  
      -pc_type none -m 200 -n 200 -log_summary
```

MatCUSPCopyTo	1	1.0	5.6108e-02
KSPGMRESOrthog	1630	1.0	5.5989e-01
KSPSolve	1	1.0	1.0202e+00

Pitfall: Repeated Host-Device Copies

- PCI-Express transfers kill performance

- Complete algorithm needs to run on device

- Problematic for explicit time-stepping, etc.

Pitfall: Wrong Data Sizes

- Data too small: Kernel launch latencies dominate

- Data too big: Out of memory

Pitfall: Function Pointers

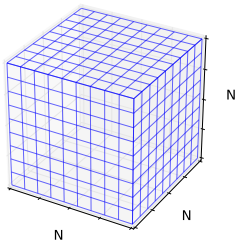
- Pass CUDA function “pointers” through library boundaries?

- OpenCL: Pass kernel sources, user-data hard to pass

- Composability?

Pitfall: GPUs are too fast for PCI-Express

Latest GPU peaks: 720 GB/sec from GPU-RAM, 16 GB/sec for PCI-Express
40x imbalance (



Compute vs. Communication

Take $N = 512$, so each field consumes 1 GB of GPU RAM

Boundary communication: $2 \times 6 \times N^2$: 31 MB

Time to load field: 1.4 ms

Time to load ghost data: **1.9 ms (!)**

**Many PETSc Operations
do NOT benefit from modern high-end GPUs
in a substantial way!**
(Exception: OpenPower systems)

Current GPU-Functionality in PETSc

	CUSP/CUDA	ViennaCL
Programming Model	CUDA	CUDA/OpenCL/OpenMP
Operations	Vector, MatMult	Vector, MatMult
Matrix Formats	CSR, ELL, HYB	CSR
Preconditioners	SA-AMG, BiCGStab	SA/Agg-AMG, Par-ILU0
MPI-related	Scatter	-

Additional Functionality

MatMult via cuSPARSE

OpenCL residual evaluation for PetscFE

PETSc on GPUs and MIC:

Current Directions

Split CUDA-buffers from CUSP

- Vector operations by cuBLAS

- MatMult by different packages

- CUSP (and others) provides add-on functionality

More CUSP Functionality in PETSc

- Relaxations (Gauss-Seidel, SOR)

- Polynomial preconditioners

- Approximate inverses

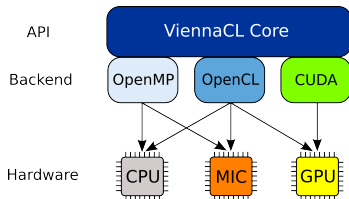
ViennaCL

CUDA, OpenCL, OpenMP backends

Backend switch at **runtime**

CUDA, OpenCL and OpenMP exposed in PETSc

Focus on shared memory machines



Recent Advances

Pipelined Krylov solvers

Fast sparse matrix-vector products

Fast sparse matrix-matrix products

Fine-grained algebraic multigrid

Fine-grained parallel ILU

Current Use of ViennaCL in PETSc

```
$> ./ex12 -vec_type viennacl -mat_type aijviennacl ...
```

Executes on OpenCL device

New Use of ViennaCL in PETSc

```
$> ./ex12 -vec_type viennacl -mat_type aijviennacl  
      -viennacl_backend openmp ...
```

Pros and Cons

- Use CPU + GPU simultaneously
- Non-intrusive, use plugin-mechanism
- Non-optimal in strong-scaling limit
- Gather experiences for best long-term solution

Currently Available

CUSP/CUDA for CUDA, ViennaCL for CUDA/OpenCL/OpenMP

Automatic use for vector operations and SpMV

Smoothed Agg. AMG via CUSP and ViennaCL

ViennaCL as CUDA/OpenCL/OpenMP-hydra

Current Activities

GPU-acceleration for GAMG

Better support for $n > 1$ processes

Use of cuBLAS and cuSPARSE



PETSc can help You

solve algebraic and DAE problems in your application area
rapidly develop efficient parallel code, can start from examples
develop new solution methods and data structures
debug and analyze performance
advice on software design, solution algorithms, and performance

`petsc-{users,dev,maint}@mcs.anl.gov`

You can help PETSc

report bugs and inconsistencies, or if you think there is a better way
tell us if the documentation is inconsistent or unclear
consider developing new algebraic methods as plugins, contribute if your
idea works