

PETSc Tutorial

GPU and Manycore CPU Support in PETSc 3.9

GPU Part

Karl Rupp (complementing Richard Mills *et al.*)

`rupp@iue.tuwien.ac.at`

`https://www.karlrupp.net/`



Institute for Microelectronics, TU Wien

Imperial College London, UK

June 4-6, 2018




*Don't believe anything
unless you can run it*

Matt Knepley

Why bother?

GFLOPs/Watt

Theoretical Peak Floating Point Operations per Watt, Double Precision



Why bother?

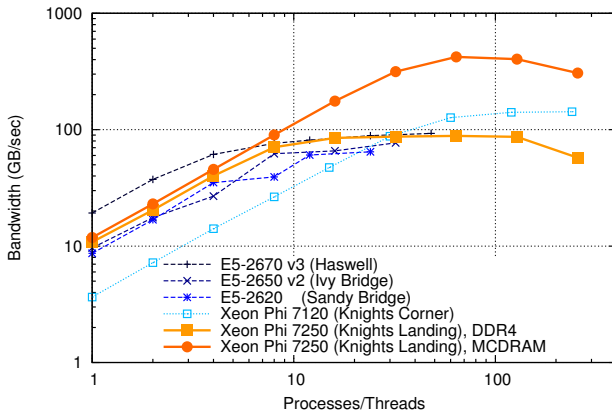
Procurements

Theta (ANL, 2016): 2nd generation INTEL Xeon Phi

Summit (ORNL, 2017), Sierra (LLNL, 2017): NVIDIA Volta GPU

~~Aurora (ANL, 2018): 3rd generation INTEL Xeon Phi~~

STREAM Benchmark Results on INTEL Hardware



PETSc on GPUs and MIC:

Current Status

Native on Xeon Phi

Forget KNC (“old Xeon Phi”)

“Just works” on KNL



CUDA

CUDA-support through CUDA/CUSPARSE

`-vec_type cuda -mat_type aijcuspars`

Only for NVIDIA GPUs



CUDA/OpenCL/OpenMP

CUDA/OpenCL/OpenMP-support through ViennaCL

`-vec_type viennacl -mat_type aijviennacl`

OpenCL on CPUs and MIC fairly poor



CUDA (CUSPARSE)

```
./configure [...] --with-cuda=1
```

Customization:

```
--with-cudac=/path/to/cuda/bin/nvcc  
--with-cuda-arch=sm_60
```

ViennaCL

```
./configure [...] --download-viennacl
```

Optional: CUDA/OpenCL/OpenMP

```
--with-cuda=1
```

```
--with-opencl-include=/path/to/OpenCL/include  
--with-opencl-lib=/path/to/libOpenCL.so
```

```
--with-openmp=1
```

How Does It Work?

Host and Device Data

```
struct _p_Vec {  
    ...  
    void          *data;           // host buffer  
    PetscCUDAFlag valid_GPU_array; // flag  
    void          *spptr;          // device buffer  
};
```

Possible Flag States

```
typedef enum {PETSC_CUDA_UNALLOCATED,  
             PETSC_CUDA_GPU,  
             PETSC_CUDA_CPU,  
             PETSC_CUDA_BOTH} PetscCUDAFlag;
```


How Does It Work?

Fallback-Operations on Host

Data becomes valid on host (PETSC_CUDA_CPU)

```
PetscErrorCode VecSetRandom_SeqCUDA_Private(..) {  
    VecGetArray(...);  
    // some operation on host memory  
    VecRestoreArray(...);  
}
```

Accelerated Operations on Device

Data becomes valid on device (PETSC_CUDA_GPU)

```
PetscErrorCode VecAYPX_SeqCUDA(..) {  
    VecCUDAGetArrayReadWrite(...);  
    // some operation on raw handles on device  
    VecCUDARestoreArrayReadWrite(...);  
}
```

KSP ex12 on Host

```
$> ./ex12  
-pc_type ilu -m 200 -n 200 -log_summary
```

```
KSPGMRESOrthog      228 1.0 6.2901e-01  
KSPSolve             1 1.0 2.7332e+00
```

KSP ex12 on Device

```
$> ./ex12 -vec_type viennacl -mat_type aijviennacl  
-pc_type ilu -m 200 -n 200 -log_summary
```

```
[0]PETSC ERROR: MatSolverPackage petsc does not support  
matrix type seqaijviennacl
```

KSP ex12 on Host

```
$> ./ex12  
      -pc_type none -m 200 -n 200 -log_summary
```

KSPGMRESOrthog	1630	1.0	4.5866e+00
KSPSolve	1	1.0	1.6361e+01


KSP ex12 on Device

```
$> ./ex12 -vec_type viennacl -mat_type aijviennacl  
      -pc_type none -m 200 -n 200 -log_summary
```

MatCUSPCopyTo	1	1.0	5.6108e-02
KSPGMRESOrthog	1630	1.0	5.5989e-01
KSPSolve	1	1.0	1.0202e+00

Pitfall 1: GPUs are too fast for PCI-Express

Latest GPU peaks: 720 GB/sec from GPU-RAM, 16 GB/sec for PCI-Express
40x imbalance (!)



Compute vs. Communication

Take $N = 512$, so each field consumes 1 GB of GPU RAM

Boundary communication: $2 \times 6 \times N^2$: 31 MB


Time to load field: 1.4 ms

Time to load ghost data: **1.9 ms (!!)**


Pitfall 2: Wrong Data Sizes

Data too small: Kernel launch latencies dominate


Data too big: Out of memory



Strong Scaling Implications



Strong Scaling Implications



Pitfall 3: Composability of GPU codes

How to pass GPU pointers through library boundaries efficiently?

High-level interfaces tend to be polluted by low-level details

**Many Non-Trivial PETSc Operations
do NOT benefit from modern high-end GPUs
in a substantial way!**
(OpenPower systems can be exceptions)

Current GPU-Functionality in PETSc

	CUDA/CUSPARSE	ViennaCL
Programming Model	CUDA	CUDA/OpenCL/OpenMP
Operations	Vector, MatMult	Vector, MatMult
Matrix Formats	CSR, ELL, HYB	CSR
Preconditioners	ILU0	SA/Agg-AMG, Par-ILU0
MPI-related	Scatter	-

Additional Functionality

MatMult via cuSPARSE

OpenCL residual evaluation for PetscFE

GPU support for SuperLU-dist

GPU support for SuiteSparse

Previous Use of ViennaCL in PETSc

```
$> ./ex12 -vec_type viennacl -mat_type aijviennacl ...
```

Executes on OpenCL device

New Use of ViennaCL in PETSc

```
$> ./ex12 -vec_type viennacl -mat_type aijviennacl  
      -viennacl_backend openmp ...
```

Pros and Cons

- Use CPU + GPU simultaneously
- Non-intrusive, use plugin-mechanism
- Non-optimal in strong-scaling limit
- Gather experiences for best long-term solution

PETSc on GPUs and MIC:

Upcoming Features for the Next PETSc Release

GPU-acceleration for GAMG

- Use GPUs on the finest levels only

- Coarse grid solvers unchanged

ViennaCL: Better support for $n > 1$ processes

- Efficient scatter operations

- Minimize PCI-Express traffic

Bindings for NVIDIA's AMGX package

- AMGX is now available as open source

- GPU-alternative to GAMG, HyPre, ML

PETSc can help You

solve algebraic and DAE problems in your application area
rapidly develop efficient parallel code, can start from examples
develop new solution methods and data structures
debug and analyze performance
advice on software design, solution algorithms, and performance

`petsc-{users,dev,maint}@mcs.anl.gov`

You can help PETSc

report bugs and inconsistencies, or if you think there is a better way
tell us if the documentation is inconsistent or unclear
consider developing new algebraic methods as plugins, contribute if your
idea works